

*Computer Vision*

---

# Covariance and Principal Components

Hansung Kim  
[h.kim@soton.ac.uk](mailto:h.kim@soton.ac.uk)

# Example



# Face dataset



# Variance and Covariance



---

# Random Variables and Expected Values

---

- ❖ Mathematicians talk variance (and covariance) in terms of *random variables* and *expected values*
- ❖ The expected value (denoted  $E[X]$ ) is the most likely value a random variable will take.
  - ❖ For this course we'll **assume** that the values an element of a feature can take are all equally likely
    - ❖ The expected value is thus just the **mean value**

---

# Variance

---

$$\sigma^2(x) = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

Variance ( $\sigma^2$ ) is the mean squared difference from the mean ( $\mu$ ).

It's a measure of how spread-out the data is.

*technically it's  $E[(X - E[X])^2]$*





---

# Covariance

---

$$\sigma(x, y) = \frac{1}{n} \sum_{i=1}^n (x - \mu_x)(y - \mu_y)$$

Covariance ( $\sigma(x, y)$ ) measures how two variables change together

*technically it's  $E[(x - E[x])(y - E[y])]$*



---

# Covariance

---

$$\sigma(x, y) = \frac{1}{n} \sum_{i=1}^n (x - \mu_x)(y - \mu_y)$$

The variance is the covariance when the two variables are the same ( $\sigma(x, x) = \sigma^2(x)$ )





---

# Covariance

---

$$\sigma(x, y) = \frac{1}{n} \sum_{i=1}^n (x - \mu_x)(y - \mu_y)$$

A covariance of 0 means the variables are uncorrelated (independent).

*Covariance is related to Correlation*

$$\rho(x, y) = \frac{\sigma(x, y)}{\sqrt{\sigma^2(x)}\sqrt{\sigma^2(y)}}$$



# Covariance Matrix

$$\Sigma = \begin{bmatrix} \sigma(X_1, X_1) & \sigma(X_1, X_2) & \dots & \sigma(X_1, X_n) \\ \sigma(X_2, X_1) & \sigma(X_2, X_2) & \dots & \sigma(X_2, X_n) \\ \vdots & \vdots & \ddots & \vdots \\ \sigma(X_n, X_1) & \sigma(X_n, X_2) & \dots & \sigma(X_n, X_n) \end{bmatrix}$$

A covariance matrix encodes how all possible pairs of dimensions in an  $n$ -dimensional dataset vary together





# Covariance Matrix







$$\Sigma = \begin{bmatrix} \sigma(X_1, X_1) & \sigma(X_1, X_2) & \dots & \sigma(X_1, X_n) \\ \sigma(X_2, X_1) & \sigma(X_2, X_2) & \dots & \sigma(X_2, X_n) \\ \vdots & \vdots & \ddots & \vdots \\ \sigma(X_n, X_1) & \sigma(X_n, X_2) & \dots & \sigma(X_n, X_n) \end{bmatrix}$$

The covariance matrix is a **square symmetric matrix**



# 2D Covariance


Geometric interpretation of  $\Sigma = \begin{bmatrix} \sigma(x, x) & \sigma(x, y) \\ \sigma(y, x) & \sigma(y, y) \end{bmatrix}$





# Mean Centring

- ❖ Mean Centring is the process of computing the mean (across each dimension independently) of a set of vectors, and then subtracting the mean vector from every vector in the set.
- ❖ All the vectors will be translated so their average position is the origin



# Covariance matrix again

$$\mathbf{Z} = \begin{bmatrix} V_{11} & V_{12} & V_{13} & \dots \\ V_{21} & V_{22} & V_{23} & \dots \\ V_{31} & V_{32} & V_{33} & \dots \\ V_{41} & V_{42} & V_{43} & \dots \end{bmatrix}$$

Each row is a **mean centred** feature vector

To the mean of  
each column

*Then*

$$\Sigma \propto \mathbf{Z}^T \mathbf{Z}$$






# Principal axes of variation

# Basis


- ❖ A basis is a set of  $n$  **linearly independent** vectors in an  $n$  dimensional space
- ❖ The vectors are **orthogonal**
- ❖ They form a “coordinate system”
- ❖ There are an infinite number of possible basis






# The first principal axis

- ❖ For a given set of  $n$  dimensional data, the *first principle axis* (or just *principal axis*) is the vector that describes the direction of **greatest variance**.



# The second principal axis


- ❖ The **second principal axis** is a vector orthogonal (perpendicular) to the first major axis.





# Three principal axes in 3D

- ❖ In a space with 3 or more dimensions, the second principal axis is the direction of the second greatest variance orthogonal to the first principal axes.
- ❖ The set of  $n$  **principal axes** of an  $n$  dimensional space are a **basis**



# Eigenvectors, Eigenvalues and Eigendecomposition



# Eigenvectors and Eigenvalues

---

Very important equation!



$$Av = \lambda v$$



# Eigenvectors and Eigenvalues

a  $n \times n$  square matrix

a scalar value,  
known as an  
**eigenvalue**

$$\boxed{A} \boxed{v} = \boxed{\lambda} \boxed{v}$$

an  $n$  dimensional vector,  
known as an **eigenvector**





# Eigenvectors and Eigenvalues

---

$$Av = \lambda v$$

There are at most  $n$  eigenvector-eigenvalue pairs

If  $A$  is **symmetric**, then the set of eigenvectors  
is **orthogonal**

*Can you see where this is going?*



# Eigenvectors and Eigenvalues

$$Av = \lambda v$$

- ❖ If  $A$  is a **covariance matrix**, then the eigenvectors are the **principal axes**.
- ❖ The eigenvalues are proportional to the **variance** of the data along each eigenvector.
- ❖ The eigenvector corresponding to the **largest eigenvalue** is the first P.C.





---

# Finding the EVecs and EVals

---

- ❖ For small matrices ( $n \leq 4$ ) there are algebraic solutions to finding all the eigenvector-eigenvalue pairs
- ❖ For larger matrices, numerical solutions to the **Eigendecomposition** must be sought.

# Eigendecomposition

columns of  $\mathbf{Q}$  are the eigenvectors

$$\mathbf{A} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1}$$

diagonal eigenvalue matrix ( $\mathbf{\Lambda}_{ii} = \lambda_i$ )





---

# Eigendecomposition

---

$$\mathbf{A} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1}$$

If  $\mathbf{A}$  is *real symmetric* (i.e. a covariance matrix), then  $\mathbf{Q}^{-1} = \mathbf{Q}^T$  (i.e. eigenvectors are orthogonal), so:

$$\mathbf{A} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T$$



---

# Eigendecomposition

---

In summary, the Eigendecomposition  
of a covariance matrix  $\mathbf{A}$ :

$$\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$$

Gives you the principal axes and  
their relative magnitudes








---

# Ordering

---

- ❖ Standard Eigendecomposition implementations will order the eigenvectors (columns of  $\mathbf{Q}$ ) such that the eigenvalues (in the diagonal of  $\mathbf{\Lambda}$ ) are sorted in order of decreasing value.
- ❖ Some solvers are optimised to only find the top  $k$  eigenvalues and corresponding eigenvectors, rather than all of them.

# Covariance, Eigendecomposition and principal axes



$$Q = \begin{bmatrix} 0.00 & 1.00 \\ 1.00 & 0.00 \end{bmatrix}$$

$$Q = \begin{bmatrix} 0.00 & 1.00 \\ 1.00 & 0.00 \end{bmatrix}$$

$$Q = \begin{bmatrix} 0.36 & -0.93 \\ 0.93 & 0.36 \end{bmatrix}$$

$$\Lambda = \begin{bmatrix} 1.00 & 0.00 \\ 0.00 & 1.00 \end{bmatrix}$$

$$\Lambda = \begin{bmatrix} 1.00 & 0.00 \\ 0.00 & 0.14 \end{bmatrix}$$


$$\Lambda = \begin{bmatrix} 1.12 & 0.00 \\ 0.00 & 0.19 \end{bmatrix}$$



# Principal Component Analysis (PCA)

# Linear Transform


- ❖ A linear transform  $W$  projects data from one space into another:
- ❖  $T = ZW$ 
  - ❖ Original data stored in the rows of  $Z$
- ❖  $T$  can have **fewer dimensions** than  $Z$ .





# Linear Transform

- ❖ The effects of a linear transform can be reversed if **W** is **invertible**:
- ❖  $Z = TW^{-1}$ 
  - ❖ A lossy process if the dimensionality of the spaces is different



---

# PCA

---

- ❖ PCA is an **Orthogonal Linear Transform** that maps data from its original space to a space defined by the principal axes of the data.
- ❖ The transform matrix  $W$  is just the eigenvector matrix  $Q$  from the Eigendecomposition of the covariance matrix of the data.
- ❖ Dimensionality reduction can be achieved by removing the eigenvectors with low eigenvalues from  $Q$  (i.e. keeping the first  $L$  columns of  $Q$  assuming the eigenvectors are sorted by decreasing eigenvalue).








# PCA Algorithm


1. Mean-centre the data vectors
2. Form the vectors into a matrix  $\mathbf{Z}$ , such that each row corresponds to a vector
3. Perform the Eigendecomposition of the matrix  $\mathbf{Z}^T\mathbf{Z}$ , to recover the eigenvector matrix  $\mathbf{Q}$  and diagonal eigenvalue matrix  $\mathbf{\Lambda}$ :  $\mathbf{Z}^T\mathbf{Z} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$
4. Sort the columns of  $\mathbf{Q}$  and corresponding diagonal values of  $\mathbf{\Lambda}$  so that the eigenvalues are decreasing.
5. Select the  $L$  largest eigenvectors of  $\mathbf{Q}$  (the first  $L$  columns) to create the transform matrix  $\mathbf{Q}_L$ .
6. Project the original vectors into a lower dimensional space,  $\mathbf{T}_L$ :  $\mathbf{T}_L = \mathbf{Z}\mathbf{Q}_L$




# PCA



3D



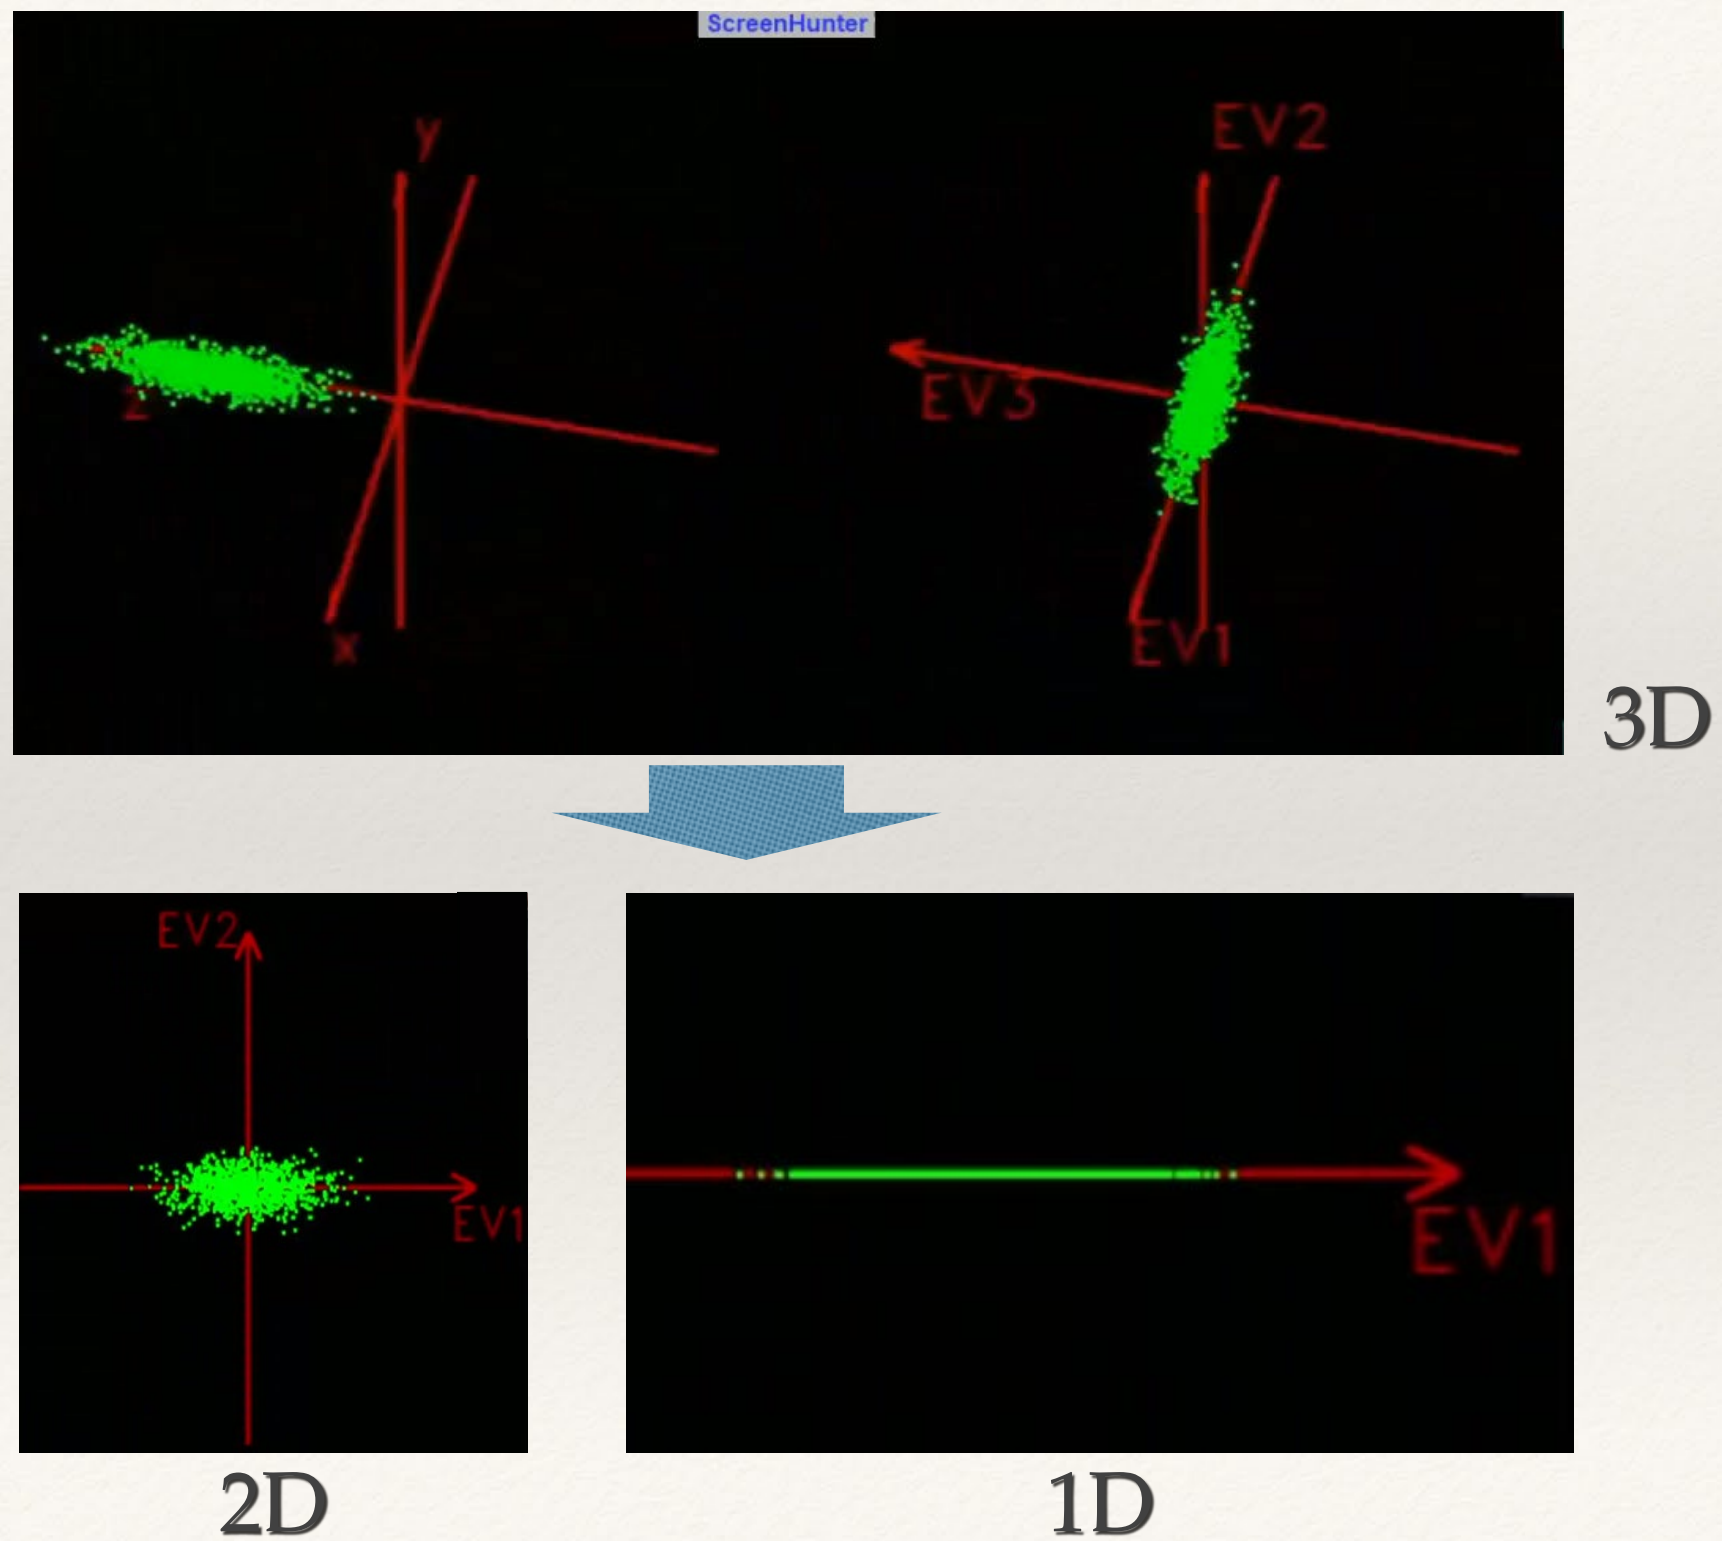
2D



1D




# PCA



# Eigenfaces (Eigenimages)



# A simple kind of feature...



---

# ... but with some problems ...

---

- ❖ Not invariant to:
  - ❖ Change in the position/orientation/scale of the object in the image
  - ❖ Changes in lighting
  - ❖ Size of the image
- ❖ Highly susceptible to image noise



---

# Making it invariant

---

- ❖ Require (almost) the same object pose across images (i.e. full frontal faces)
- ❖ Align (rotate, scale and translate) the images so that the a common feature is in the same place (i.e. the eyes in a set of face images)
- ❖ Make all the aligned images the same size
- ❖ (optional) Normalise (or perhaps histogram equalise) the images so they are invariant to global intensity changes

---

# ...but there is still a bit of a problem...

---

- ❖ The feature vectors are huge!
  - ❖ If the images are 100x200 pixels, the vector has 20000 dimensions
  - ❖ That's not really practical...
  - ❖ Also, the vectors are still highly susceptible to imaging noise and variations due to slight mis-alignments



---

# Potential solution... Apply PCA

---

- ❖ PCA can be used to reduce the dimensionality
  - ❖ smaller number of dimensions allows greater robustness to noise and mis-alignment
    - ❖ there are fewer degrees of freedom, so noise/mis-alignment has much less effect
    - ❖ and the dominant features are captured
- ❖ Fewer dimensions makes applying machine-learning much more tractable.

# EigenFaces and reconstruction



# Face dataset



---

# Mean-face

---






# Mean-centred faces





# Example for two components

- ❖ Visualisation of the face datapoints projected on the first two principal components





# Reconstructed faces



Org

DIM:10

DIM:62

DIM:102

DIM:250

---

# Summary

---

- ❖ Covariance measures the “shape” of data by measuring how different dimensions change together.
- ❖ The principle axes are a basis, aligned such they describe most directions of greatest variance.
- ❖ The Eigendecomposition of the covariance matrix produces pairs of eigenvectors (corresponding to the principal axes) and eigenvalues (proportional to the variance along the respective axis).
- ❖ PCA aligns data with its principal axes, and allows dimensionally reduction by discounting axes with low variance.
- ❖ Eigenfaces applies PCA to vectors made from pixel values to make robust low-dimensional image descriptors.



# Further reading and exercises

## ❖ Further reading

- ❖ Mark's book covers PCA in the appendices.
- ❖ Wikipedia has good coverage of all the key ideas:
  - ❖ <http://en.wikipedia.org/wiki/Variance>
  - ❖ <http://en.wikipedia.org/wiki/Covariance>
  - ❖ [http://en.wikipedia.org/wiki/Covariance matrix](http://en.wikipedia.org/wiki/Covariance_matrix)
  - ❖ <http://en.wikipedia.org/wiki/Eigenvalue>, [eigenvector](http://en.wikipedia.org/wiki/Eigenvector) and [eigenspace](http://en.wikipedia.org/wiki/Eigenspace)
  - ❖ [http://en.wikipedia.org/wiki/Eigendecomposition of a matrix](http://en.wikipedia.org/wiki/Eigendecomposition_of_a_matrix)
  - ❖ <http://en.wikipedia.org/wiki/Eigenface>

## ❖ Practical exercises – EigenFaces and PCA

- ❖ OpenIMAJ tutorial chapter 13
- ❖ <https://sandipanweb.wordpress.com/2018/01/06/eigenfaces-and-a-simple-face-detector-with-pca-svd-in-python/>
- ❖ <https://learnopencv.com/eigenface-using-opencv-c-python/>