



Computer Vision


Local interest points

Hansung Kim
h.kim@soton.ac.uk

- ❖ Finding stable (repeatable) interest points is a key problem in modern computer vision
 - ❖ Applications in areas such as tracking, matching, image alignment, making robust features for classification and search, robot navigation, 3D reconstruction, ...
 - ❖ We'll look at some of these in more detail in future lectures

What makes a good interest point?

- ❖ Invariance to brightness change (local changes as well as global ones)
- ❖ Sufficient texture variation in the local neighbourhood
 - ❖ But not too much!
- ❖ Invariance to changes between the angle/position of the scene to the camera




Stable local interest points




So, how do we find them?

- ❖ Lots of different types of interest point types to choose from.
- ❖ We'll focus on two specific types and look in detail at common detection algorithms:
 - ❖ Corner detection - *Harris and Stephens*
 - ❖ Blob Detection - *Difference-of-Gaussian Extrema*

A blob is a region of an image in which some properties are constant or approximately constant -Wikipedia




corners



blobs




The Harris and Stephens corner detector




Basic Idea


- ❖ Search for corners by looking through a small window.
- ❖ Shifting that window by a small amount in *any direction* should give a *large change* in intensity




Basic Idea



“flat” region: no
change in all
directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions



Harris & Stephens: Mathematics


Weighted average change in intensity between a window and a shifted version [by $(\Delta x, \Delta y)$] of that window:

$$E(x, y) = \sum_w f(x_i, y_i) [I(x_i, y_i) - I(x_i + \Delta x, y_i + \Delta y)]^2$$


weighting
function

intensity
in window

intensity in
shifted window



flat



Gaussian



Harris & Stephens: Mathematics

- ❖ The Taylor expansion allows us to approximate the shifted intensity.
- ❖ Taking the first order terms we get this:

$$I(x_i + \Delta x, y_i + \Delta y) \approx I(x_i, y_i) + [I_x(x_i, y_i) \quad I_y(x_i, y_i)] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

(partial derivatives of the image)



Harris & Stephens: Mathematics

- ❖ Substituting and simplifying gives:

$$\begin{aligned} E(x, y) &= \sum_w [I(x_i, y_i) - I(x_i + \Delta x, y_i + \Delta y)]^2 \\ &= \sum_w \left(I(x_i, y_i) - I(x_i, y_i) - [I_x(x_i, y_i) \quad I_y(x_i, y_i)] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= \sum_w \left(-[I_x(x_i, y_i) \quad I_y(x_i, y_i)] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= \sum_w \left([I_x(x_i, y_i) \quad I_y(x_i, y_i)] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= [\Delta x \quad \Delta y] \begin{bmatrix} \sum_w (I_x(x_i, y_i))^2 & \sum_w I_x(x_i, y_i) I_y(x_i, y_i) \\ \sum_w I_x(x_i, y_i) I_y(x_i, y_i) & \sum_w (I_y(x_i, y_i))^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \\ &= [\Delta x \quad \Delta y] \mathbf{M} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \end{aligned}$$



Structure Tensor

The **square symmetric** matrix **M** is called the *Structure Tensor* or the *Second Moment matrix*

$$\mathbf{M} = \begin{bmatrix} \sum_w (I_x(x_i, y_i))^2 & \sum_w I_x(x_i, y_i) I_y(x_i, y_i) \\ \sum_w I_x(x_i, y_i) I_y(x_i, y_i) & \sum_w (I_y(x_i, y_i))^2 \end{bmatrix}$$

It concisely encodes the how the local shape intensity function of the window changes with small shifts


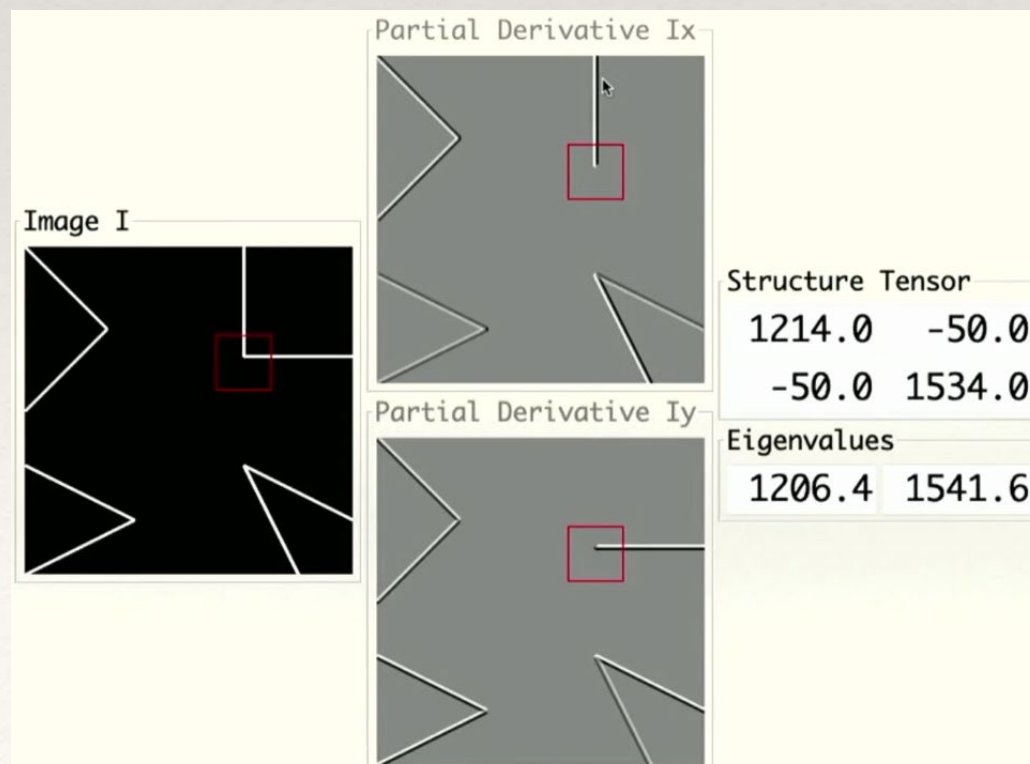

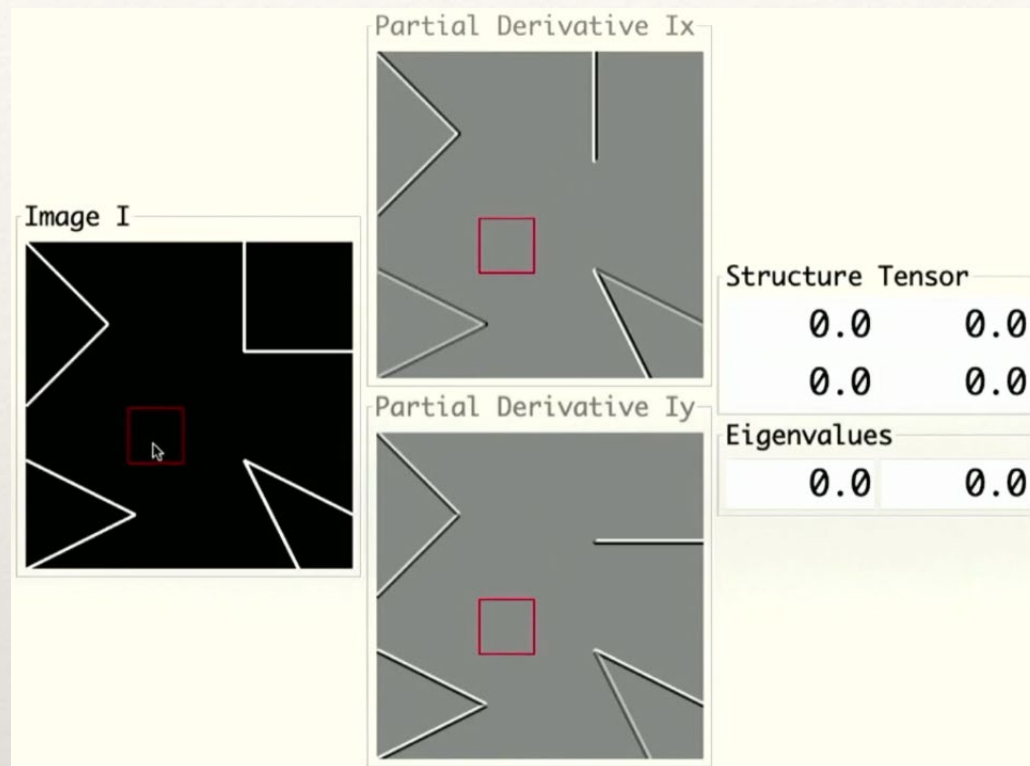


Eigenvalues of the Structure Tensor

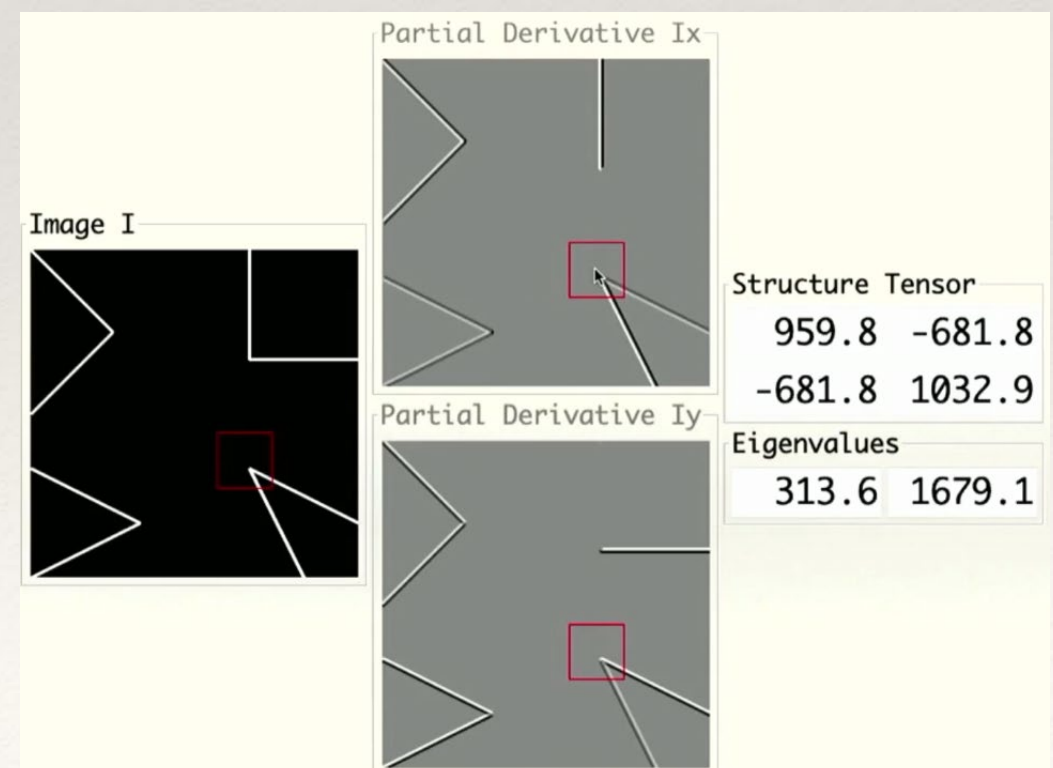
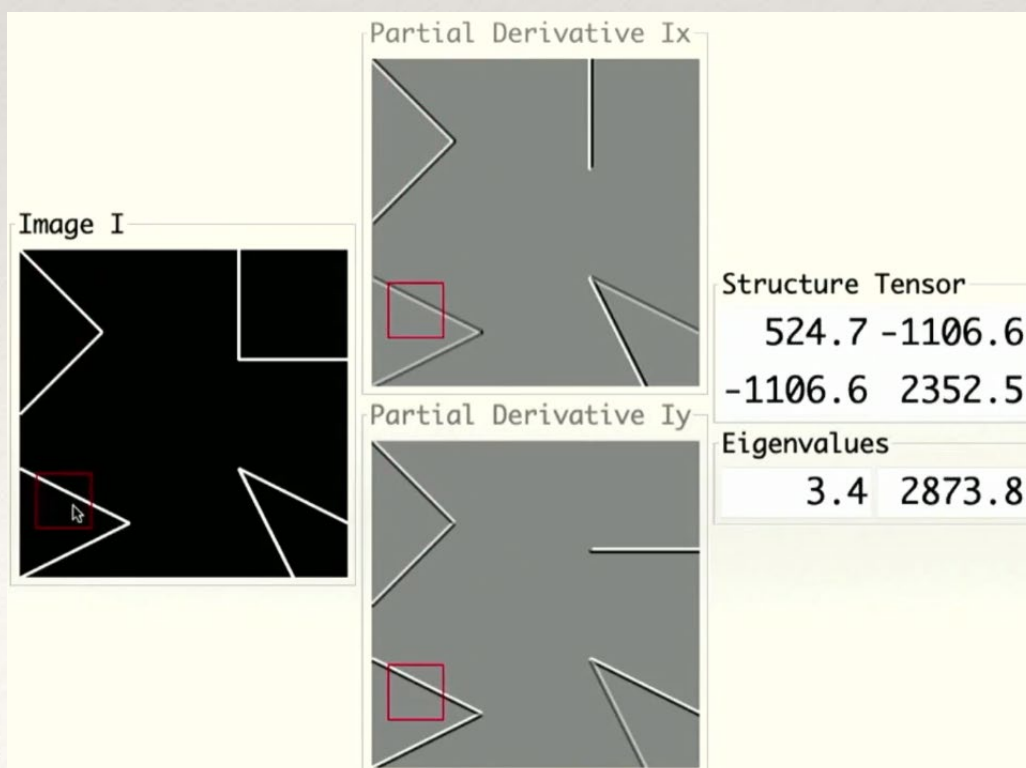
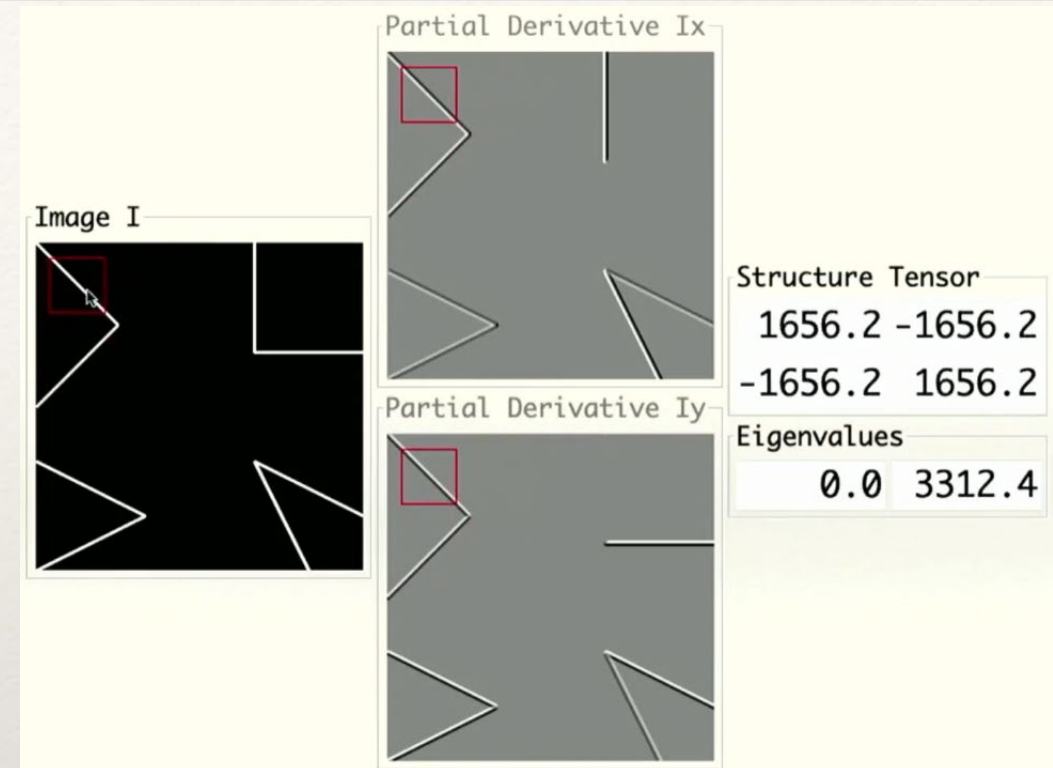

- ❖ Think back to Lecture 12 where we looked at covariance matrices...
- ❖ The eigenvalues and vectors tell us the rates of change and their respective directions




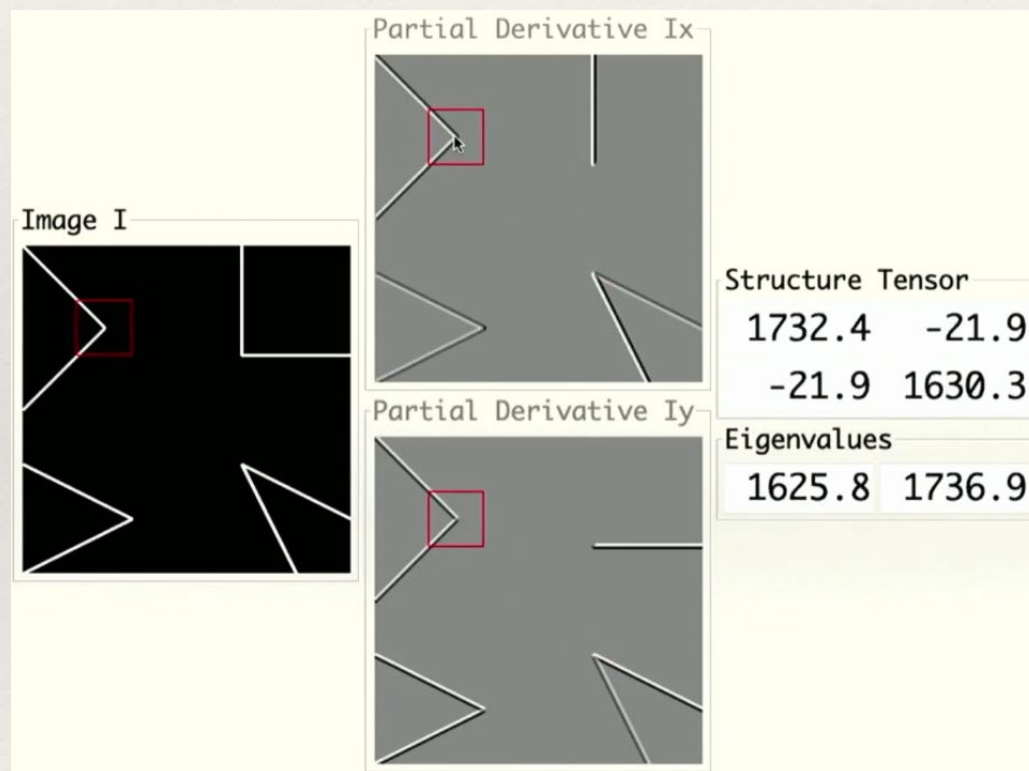
Eigenvalues of the Structure Tensor




Eigenvalues of the Structure Tensor



Eigenvalues of the Structure Tensor





Harris & Stephens Response Function

- ❖ Rather than compute the eigenvalues directly, Harris and Stephens defined a corner response function in terms of the determinant and trace of \mathbf{M} :

$$\det(\mathbf{M}) = M_{00}M_{11} - M_{01}M_{10} = M_{00}M_{11} - M_{10}^2 = \lambda_1\lambda_2$$


$$\text{trace}(\mathbf{M}) = M_{00} + M_{11} = \lambda_1 + \lambda_2$$

$$R = \det(\mathbf{M}) - k \text{trace}(\mathbf{M})^2$$

k is a small empirically set
constant (usually 0.04 - 0.06)



Harris & Stephens Response Function





Harris & Stephens Detector

- ❖ Simple algorithm:
 - ❖ Take all points with the response value above a threshold
 - ❖ Then keep only the points that are local maxima (i.e. where the current response is bigger than the 8 neighbouring pixels)



Harris & Stephens Detector



Scale in Computer Vision

The problem of scale

- ❖ As an object moves closer to the camera it get larger with more detail... as it moves further away it gets smaller and loses detail...
- ❖ If you're using a technique that uses a fixed size processing window (e.g. Harris corners, or indeed anything that involves a fixed kernel) then this is a bit of a problem!



Scale space theory

- ❖ Scale space theory is a formal framework for handling the scale problem.
- ❖ Represents the image by a series of increasingly smoothed/blurred images parameterised by a scale parameter t .
- ❖ t represents the amount of smoothing.
- ❖ **Key notion:** Image structures smaller than \sqrt{t} have been smoothed away at scale t .



The Gaussian Scale Space

- ❖ Many types of scale space are possible (depending on the smoothing function), but only the Gaussian function has the desired properties for image representation.
- ❖ These provable properties are called the “*scale space axioms*”.

Gaussian Scale Space

Formally, Gaussian scale space defined as:

$$L(\cdot, \cdot; t) = g(\cdot, \cdot; t) * f(\cdot, \cdot)$$

Note: convolution is over x, y for fixed t

where $t \geq 0$ and,

$$g(x, y; t) = \frac{1}{2\pi t} e^{-(x^2 + y^2)/2t}$$

Note: $t = \sigma^2$ = variance of the Gaussian



Gaussian Scale Space

Normally, only a fixed set of values of t are used
- it's common to use integer powers of 2 or $\sqrt{2}$



$t=0$




$t=1$




$t=2$



$t=4$



$t=16$




$t=32$



Nyquist-Shannon Sampling theorem

If a function $x(t)$ contains no frequencies higher than B hertz, it is completely determined by giving its ordinates at a series of points spaced $1/(2B)$ seconds apart.




...so, if you filter the signal with a low-pass filter that halves the frequency content, you can also half the sampling rate without loss of information...



Gaussian Pyramid

- ❖ Every time you double t in scale space, you can half the image size without loss of information!
- ❖ Leads to a much more efficient representation
 - ❖ faster processing
 - ❖ less memory




Multi-scale Harris & Stephens

- ❖ Extending the Harris and Stephens detector to work across scales is easy...
- ❖ We define a Gaussian scale space with a fixed set of scales and compute the corner response function at every pixel of each scale and keep only those with a response above a certain threshold.



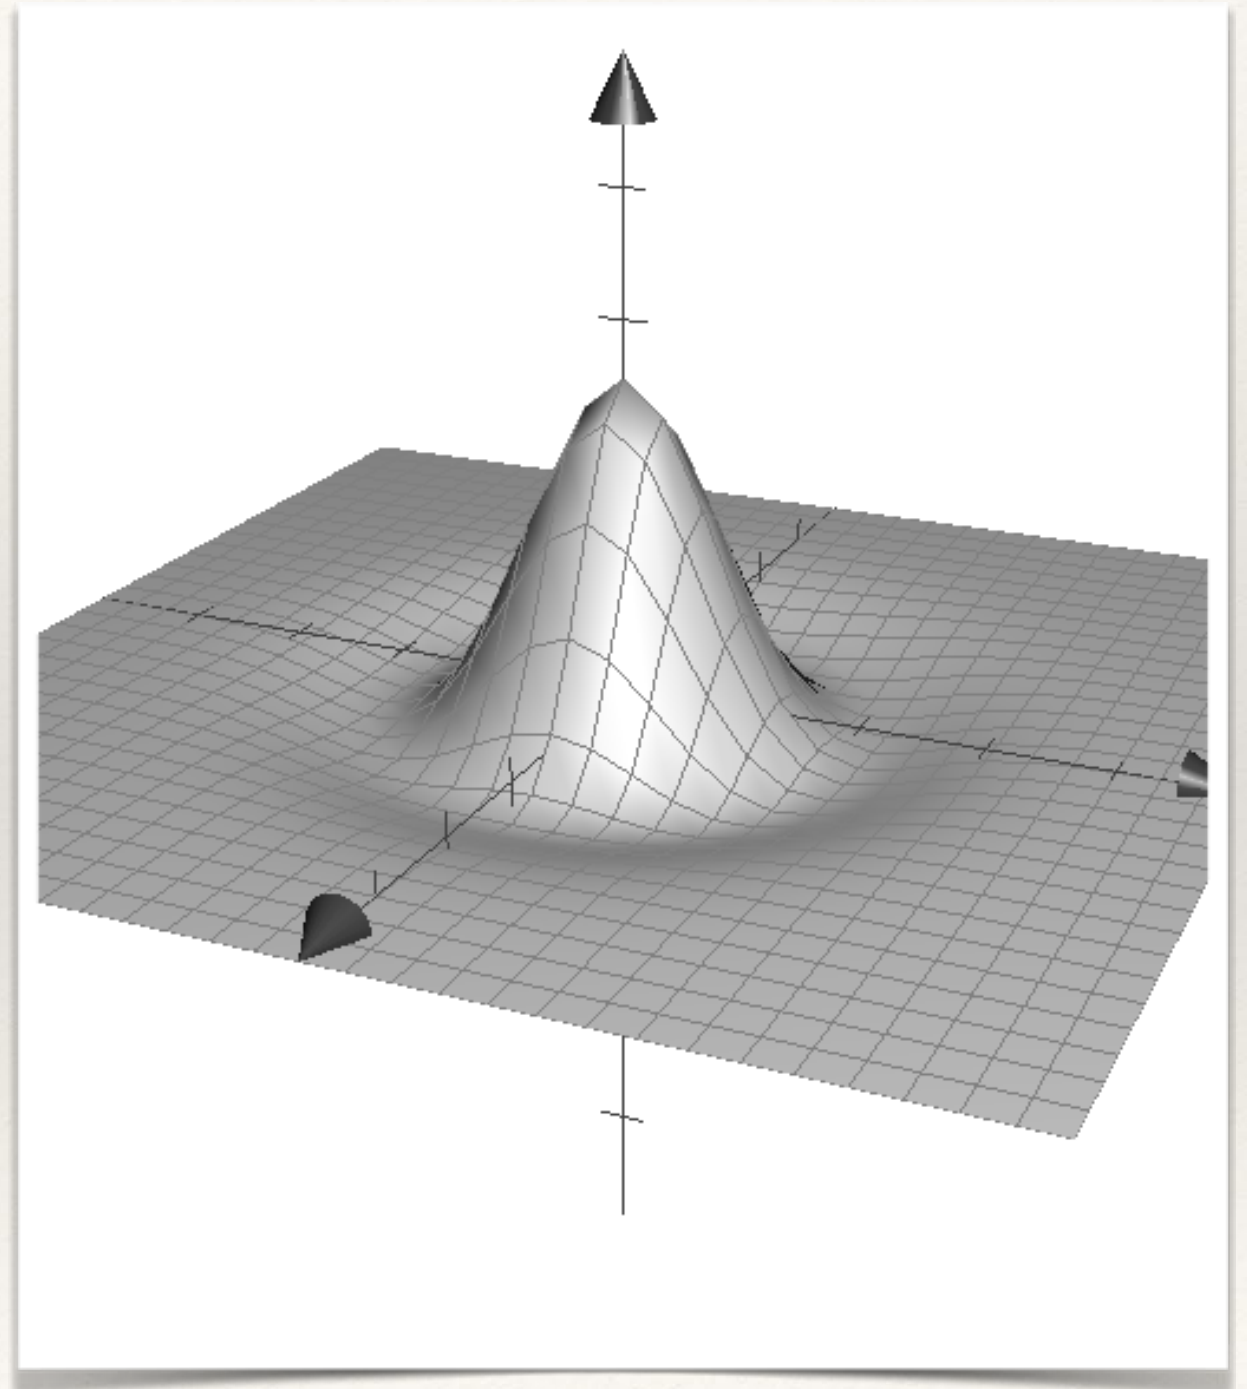
Multi-scale Harris & Stephens Corner Detector



Blob Detection

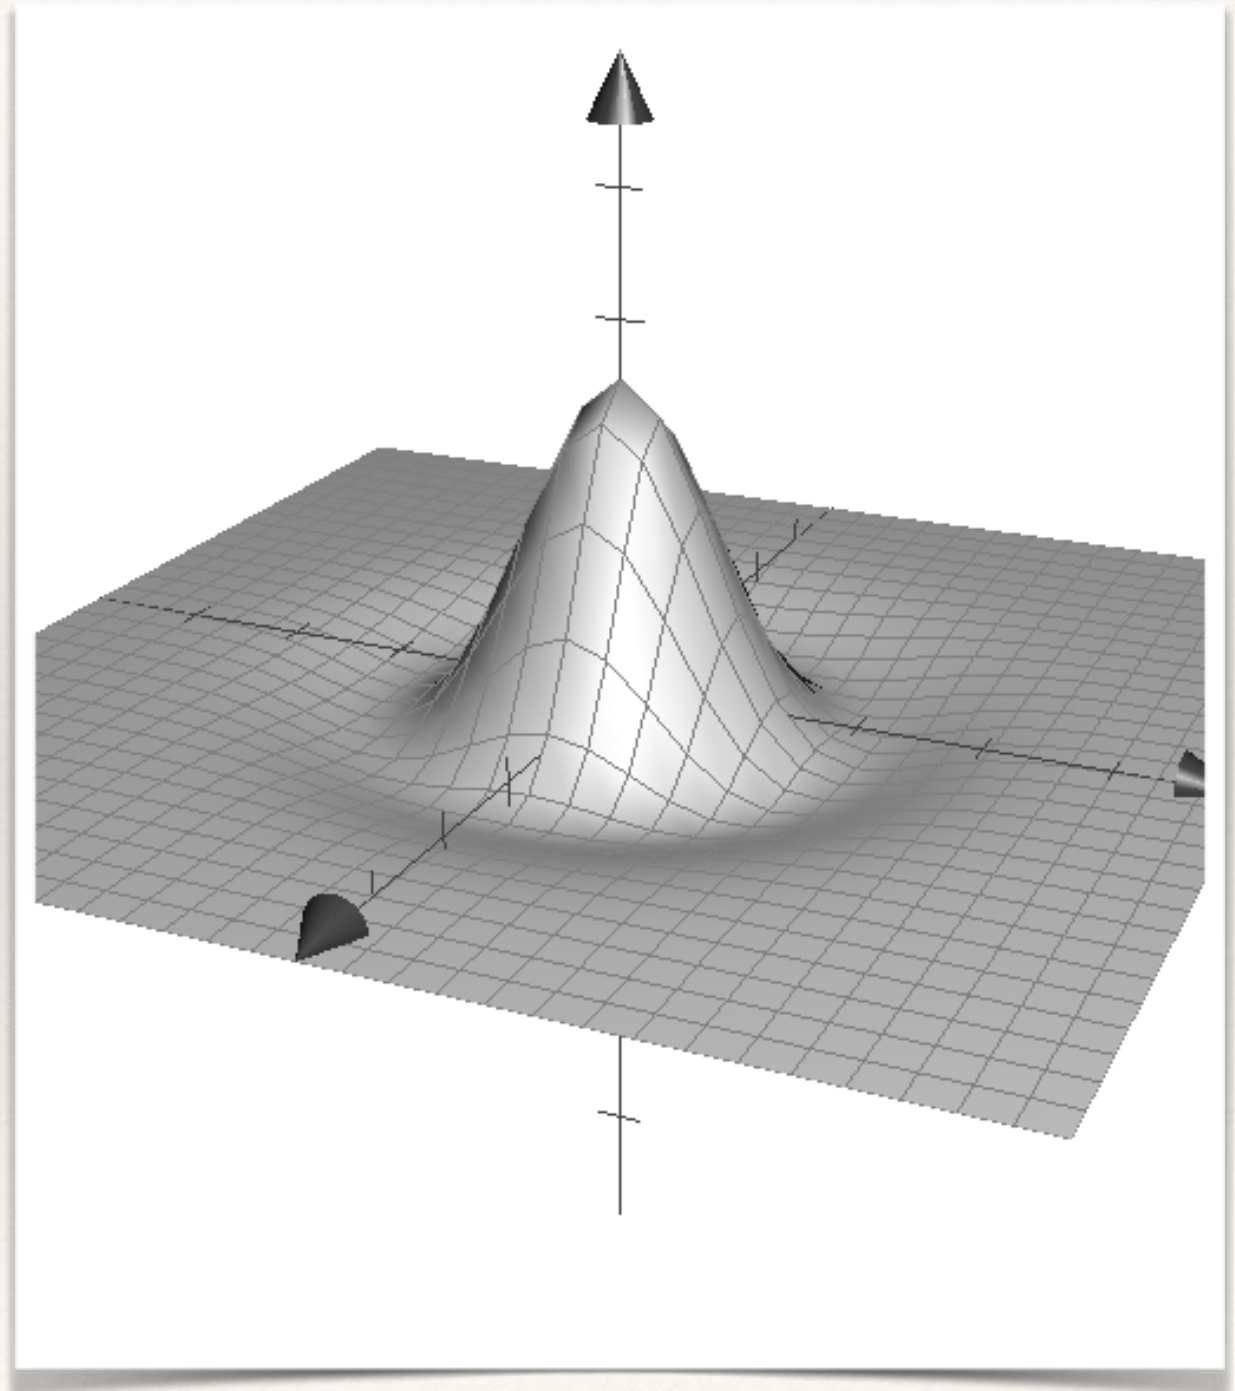
Recap: Laplacian of Gaussian

- ❖ Recall that the LoG is the second derivative of a Gaussian
- ❖ Used in the Marr-Hildreth edge detector
 - ❖ Zero crossings of LoG convolution




Laplacian of Gaussian

- ❖ Instead of finding zero crossings, find local minima or maxima => blob detector!



Scale space LoG

- ❖ Normalised scale space LoG defined as:
$$\nabla_{norm}^2 L(x, y; t) = t(L_{xx} + L_{yy})$$
- ❖ By finding extrema of this function in scale space, you can find *blobs* at their representative scale ($\sim \sqrt{2t}$)
- ❖ Just need to look at the neighbouring pixels!



Scale space DoG

- ❖ In practice it's computationally expensive to build a LoG scale space.

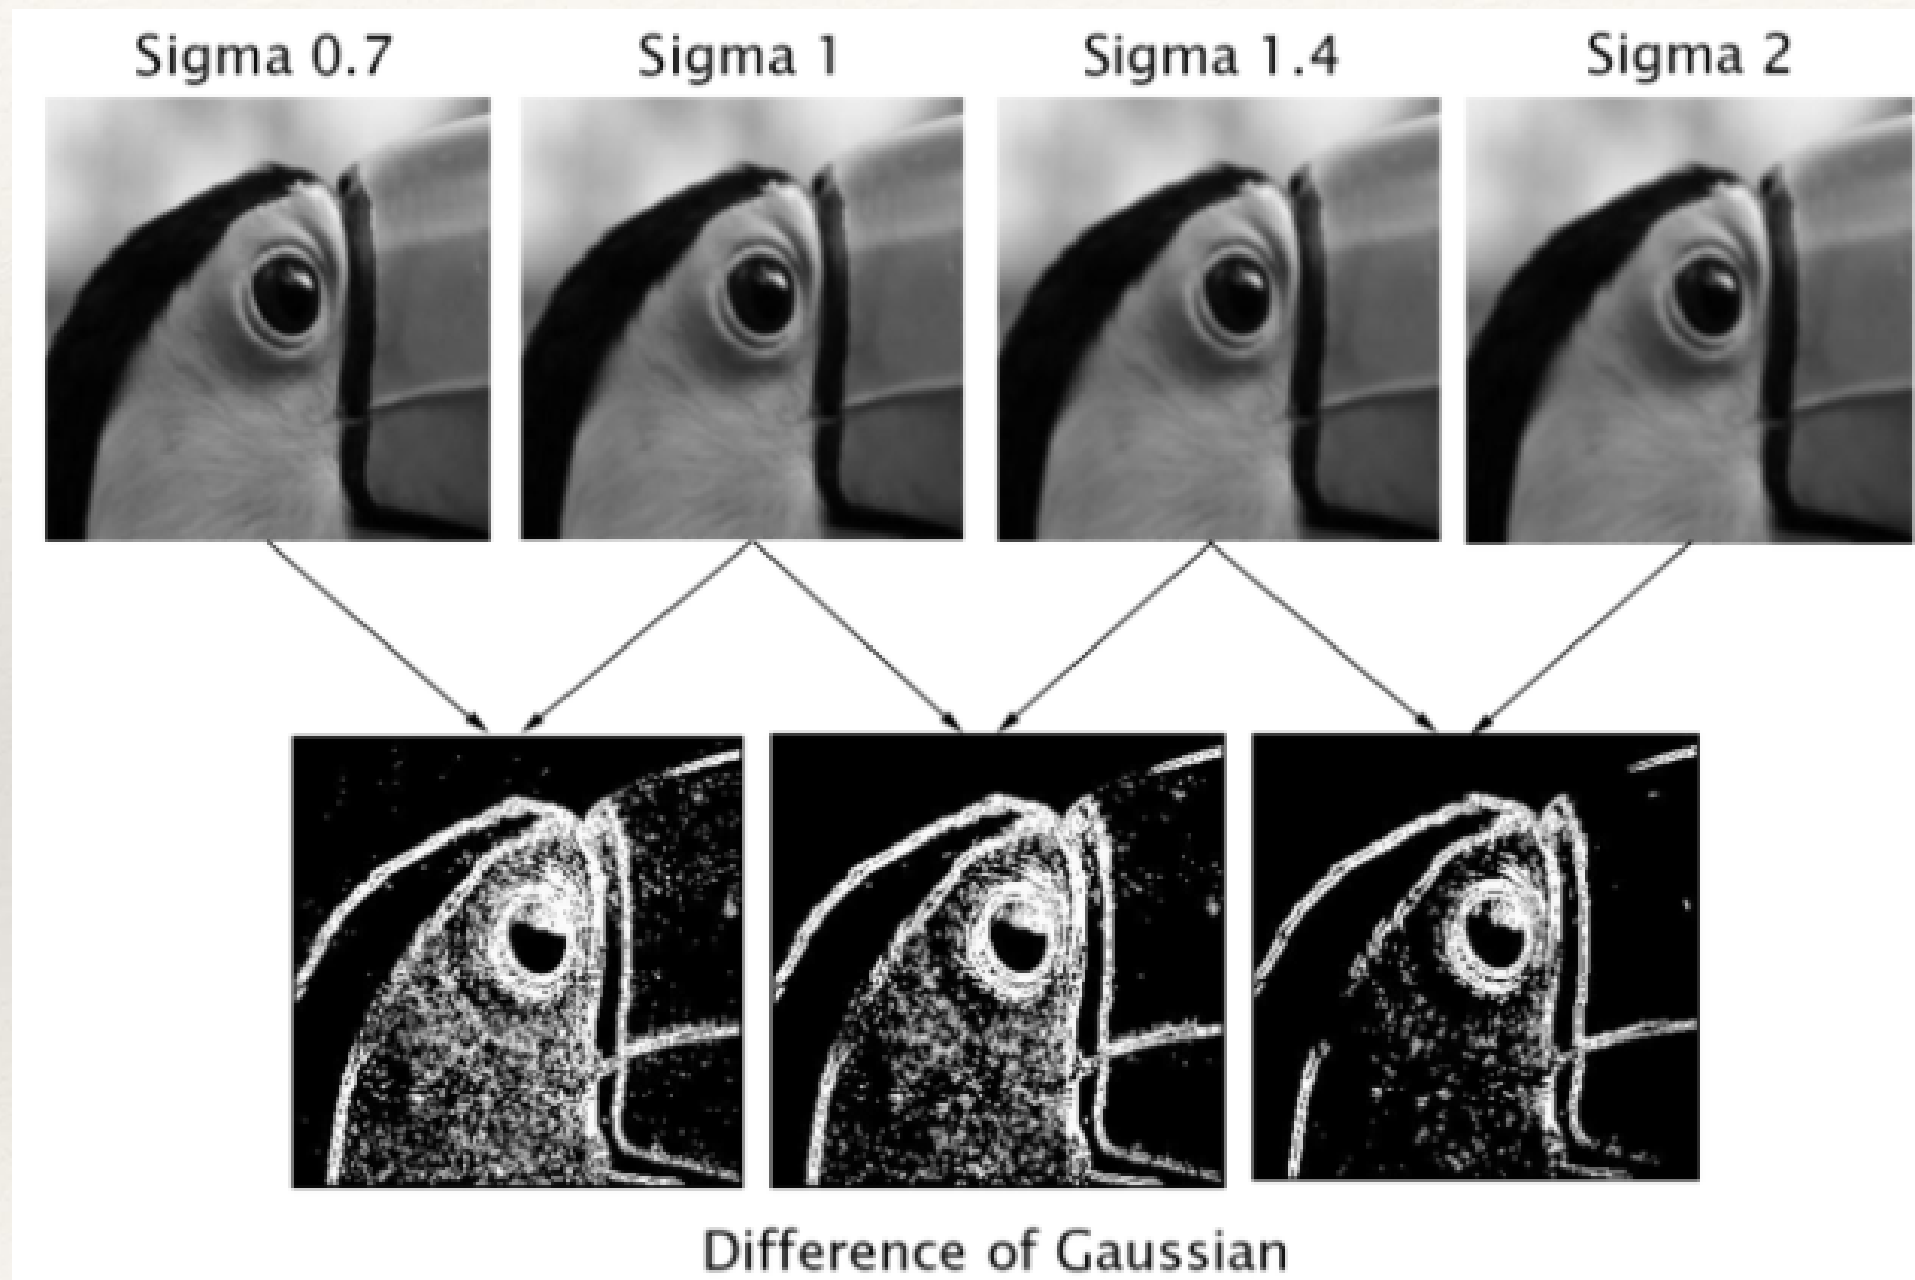
- ❖ But, the following approximation can be made:

$$\nabla_{norm}^2 L(x, y; t) \approx \frac{t}{\Delta t} (L(x, y; t + \Delta t) - L(x, y; t - \Delta t))$$

- ❖ This is called a Difference-of-Gaussians (*DoG*)
 - ❖ Implies that the LoG scale space can be built from subtracting adjacent scales of a Gaussian scale space



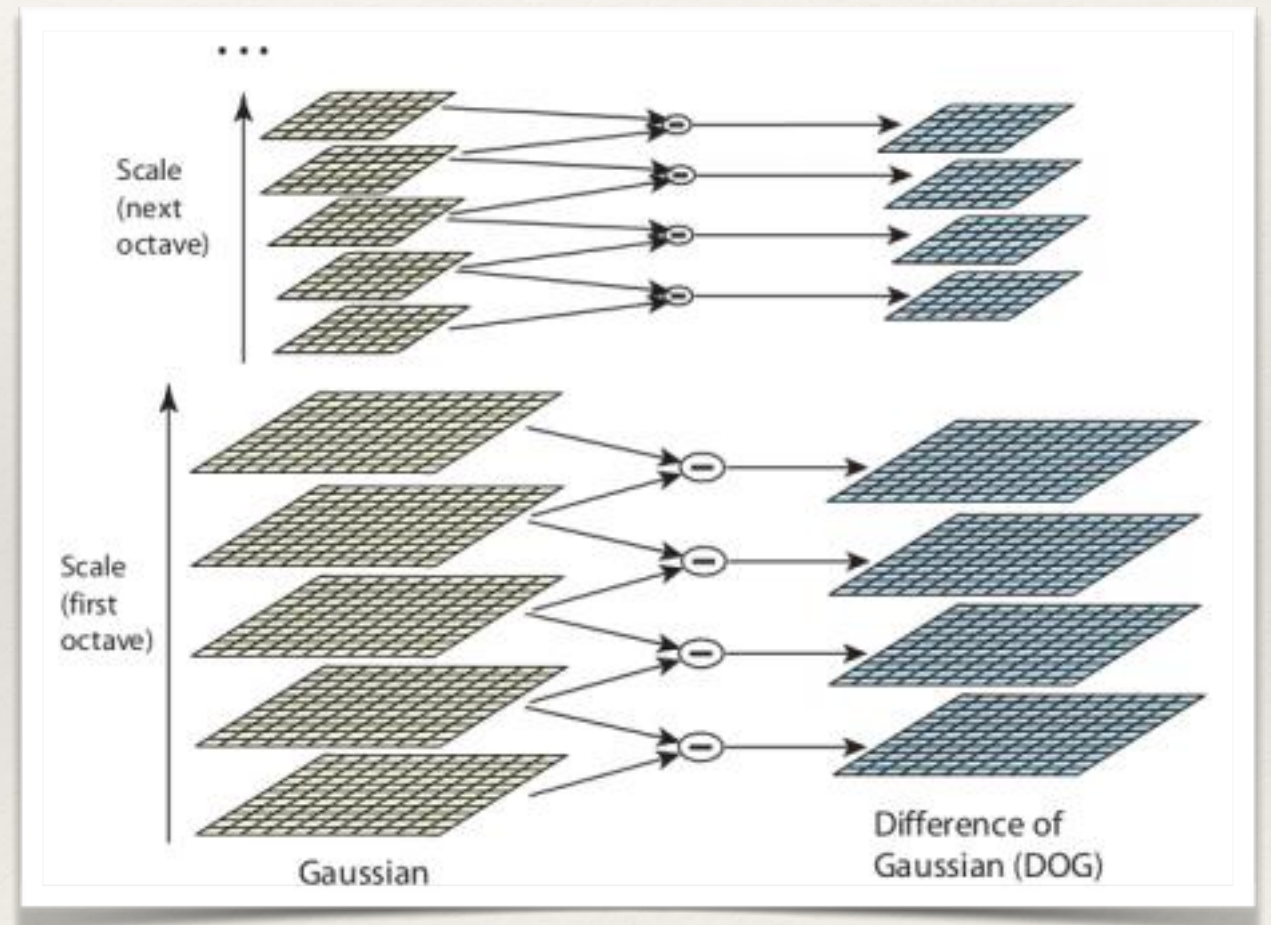
Difference of Gaussian Response





Very useful property: *if a blob is detected at $(x_0, y_0; t_0)$ in an image, then under a scaling of that image by a factor s , the same blob would be detected at $(sx_0, sy_0; s^2t_0)$ in the scaled image.*

DoG Pyramid

- ❖ Of course, for efficiency you can also build a DoG pyramid
- ❖ an *oversampled* pyramid as there are multiple images between a doubling of scale.
- ❖ Images between a doubling of scale are an *octave*.



Multi-scale Difference of Gaussian Blob detector



Summary

- ❖ Interest points have loads of applications in computer vision.
 - ❖ They need to be robustly detected, and invariant to rotation, lighting change, etc.
- ❖ We've looked at two types: corners and blobs
 - ❖ Harris & Stephens is a common corner detector
 - ❖ Finding extrema in a multi scale DoG pyramid provides a robust blob detector
- ❖ Scale space theory allows us to find features (corners and blobs) of different sizes

Further reading

- ❖ Harris & Stephens
 - ❖ Original paper by C. Harris and M. Stephens (1988)
<http://www.bmva.org/bmvc/1988/avc-88-023.pdf>
 - ❖ Material also partially covered in Mark's book pp. 159-165
- ❖ Scale space http://en.wikipedia.org/wiki/Scale_space_representation
- ❖ Blob detection
 - ❖ Wikipedia
[http://en.wikipedia.org/wiki/Blob_detection#The Laplacian of Gaussian](http://en.wikipedia.org/wiki/Blob_detection#The_Laplacian_of_Gaussian)
 - ❖ https://medium.com/@vad710/cv-for-busy-devs-improving-features-df20c3aa5887#_
 - ❖ Lowe's seminal SIFT paper also has a good description of finding blobs using the DoG in section 3:
<http://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

Exercises

❖ Practical exercises

❖ Harris Corner Detection:

- ❖ `org.openimaj.image.feature.local.interest.HarrisIPD` class in OpenIMAJ

- ❖ https://opencv24-python-tutorials.readthedocs.io/en/latest/py_tutorials/py_feature2d/py_features_harris/py_features_harris.html

❖ DoG

- ❖ Dog is a part of SIFT in OpenCV

- ❖ https://opencv24-python-tutorials.readthedocs.io/en/latest/py_tutorials/py_feature2d/py_sift_intro/py_sift_intro.html

- ❖ Chapter 5 of the OpenIMAJ tutorial covers finding DoG blobs