

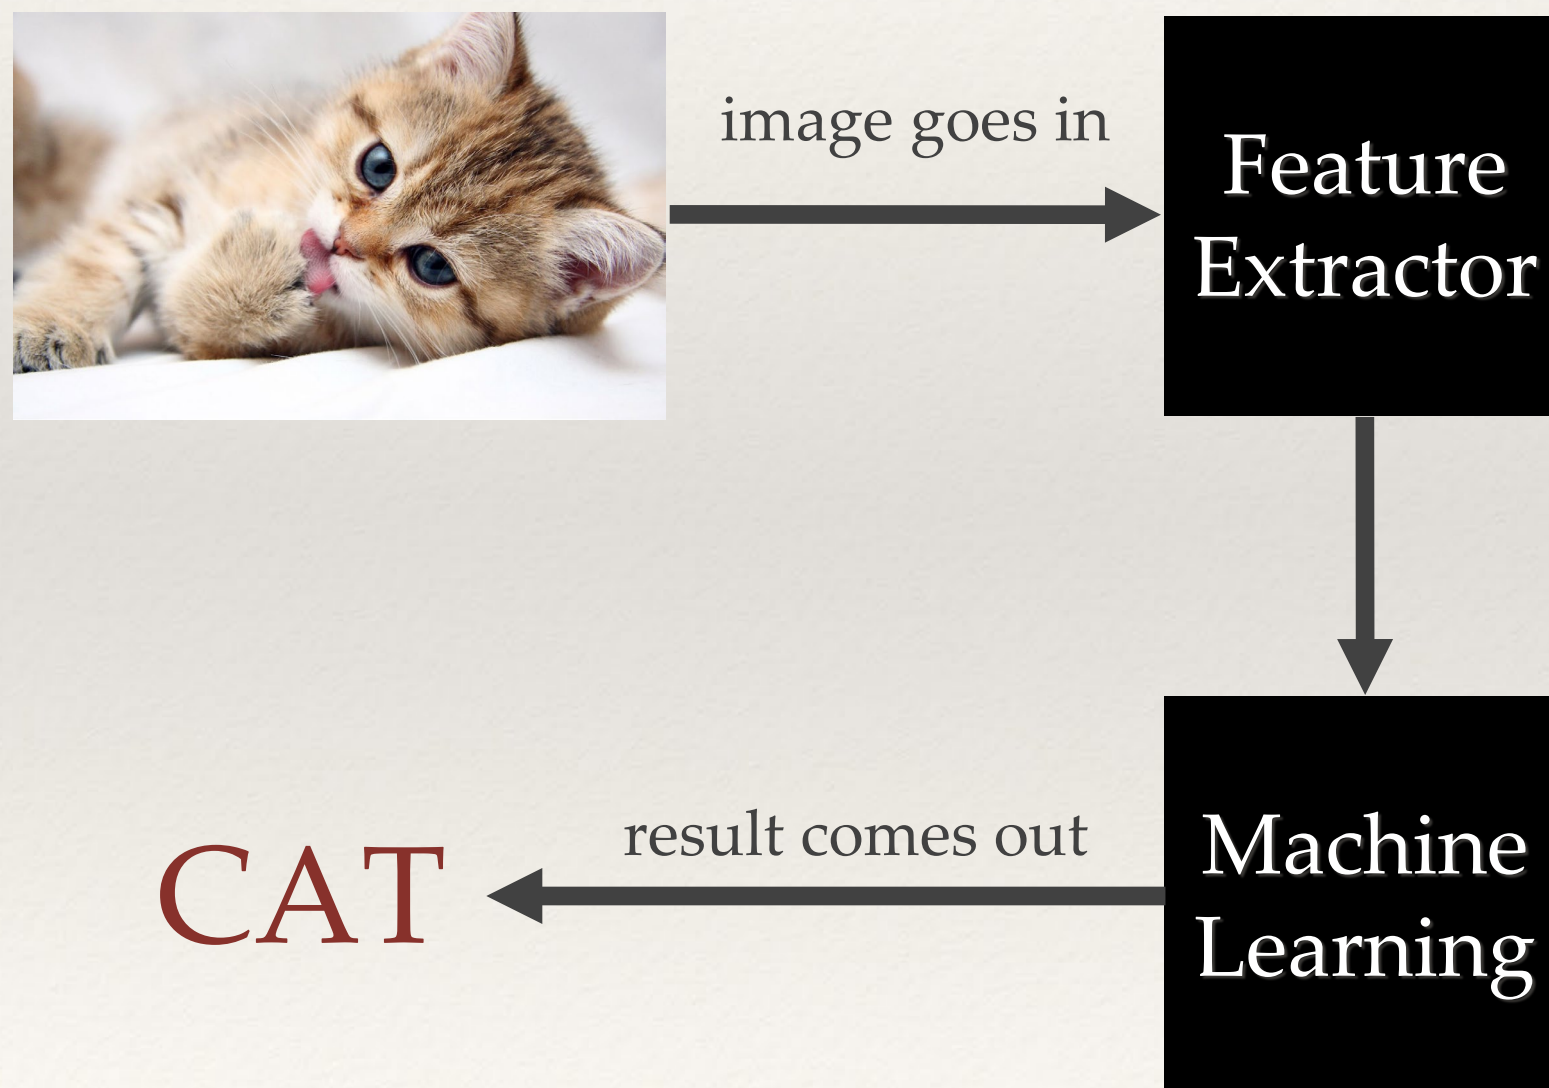
*Computer Vision*

# Image classification and auto-annotation

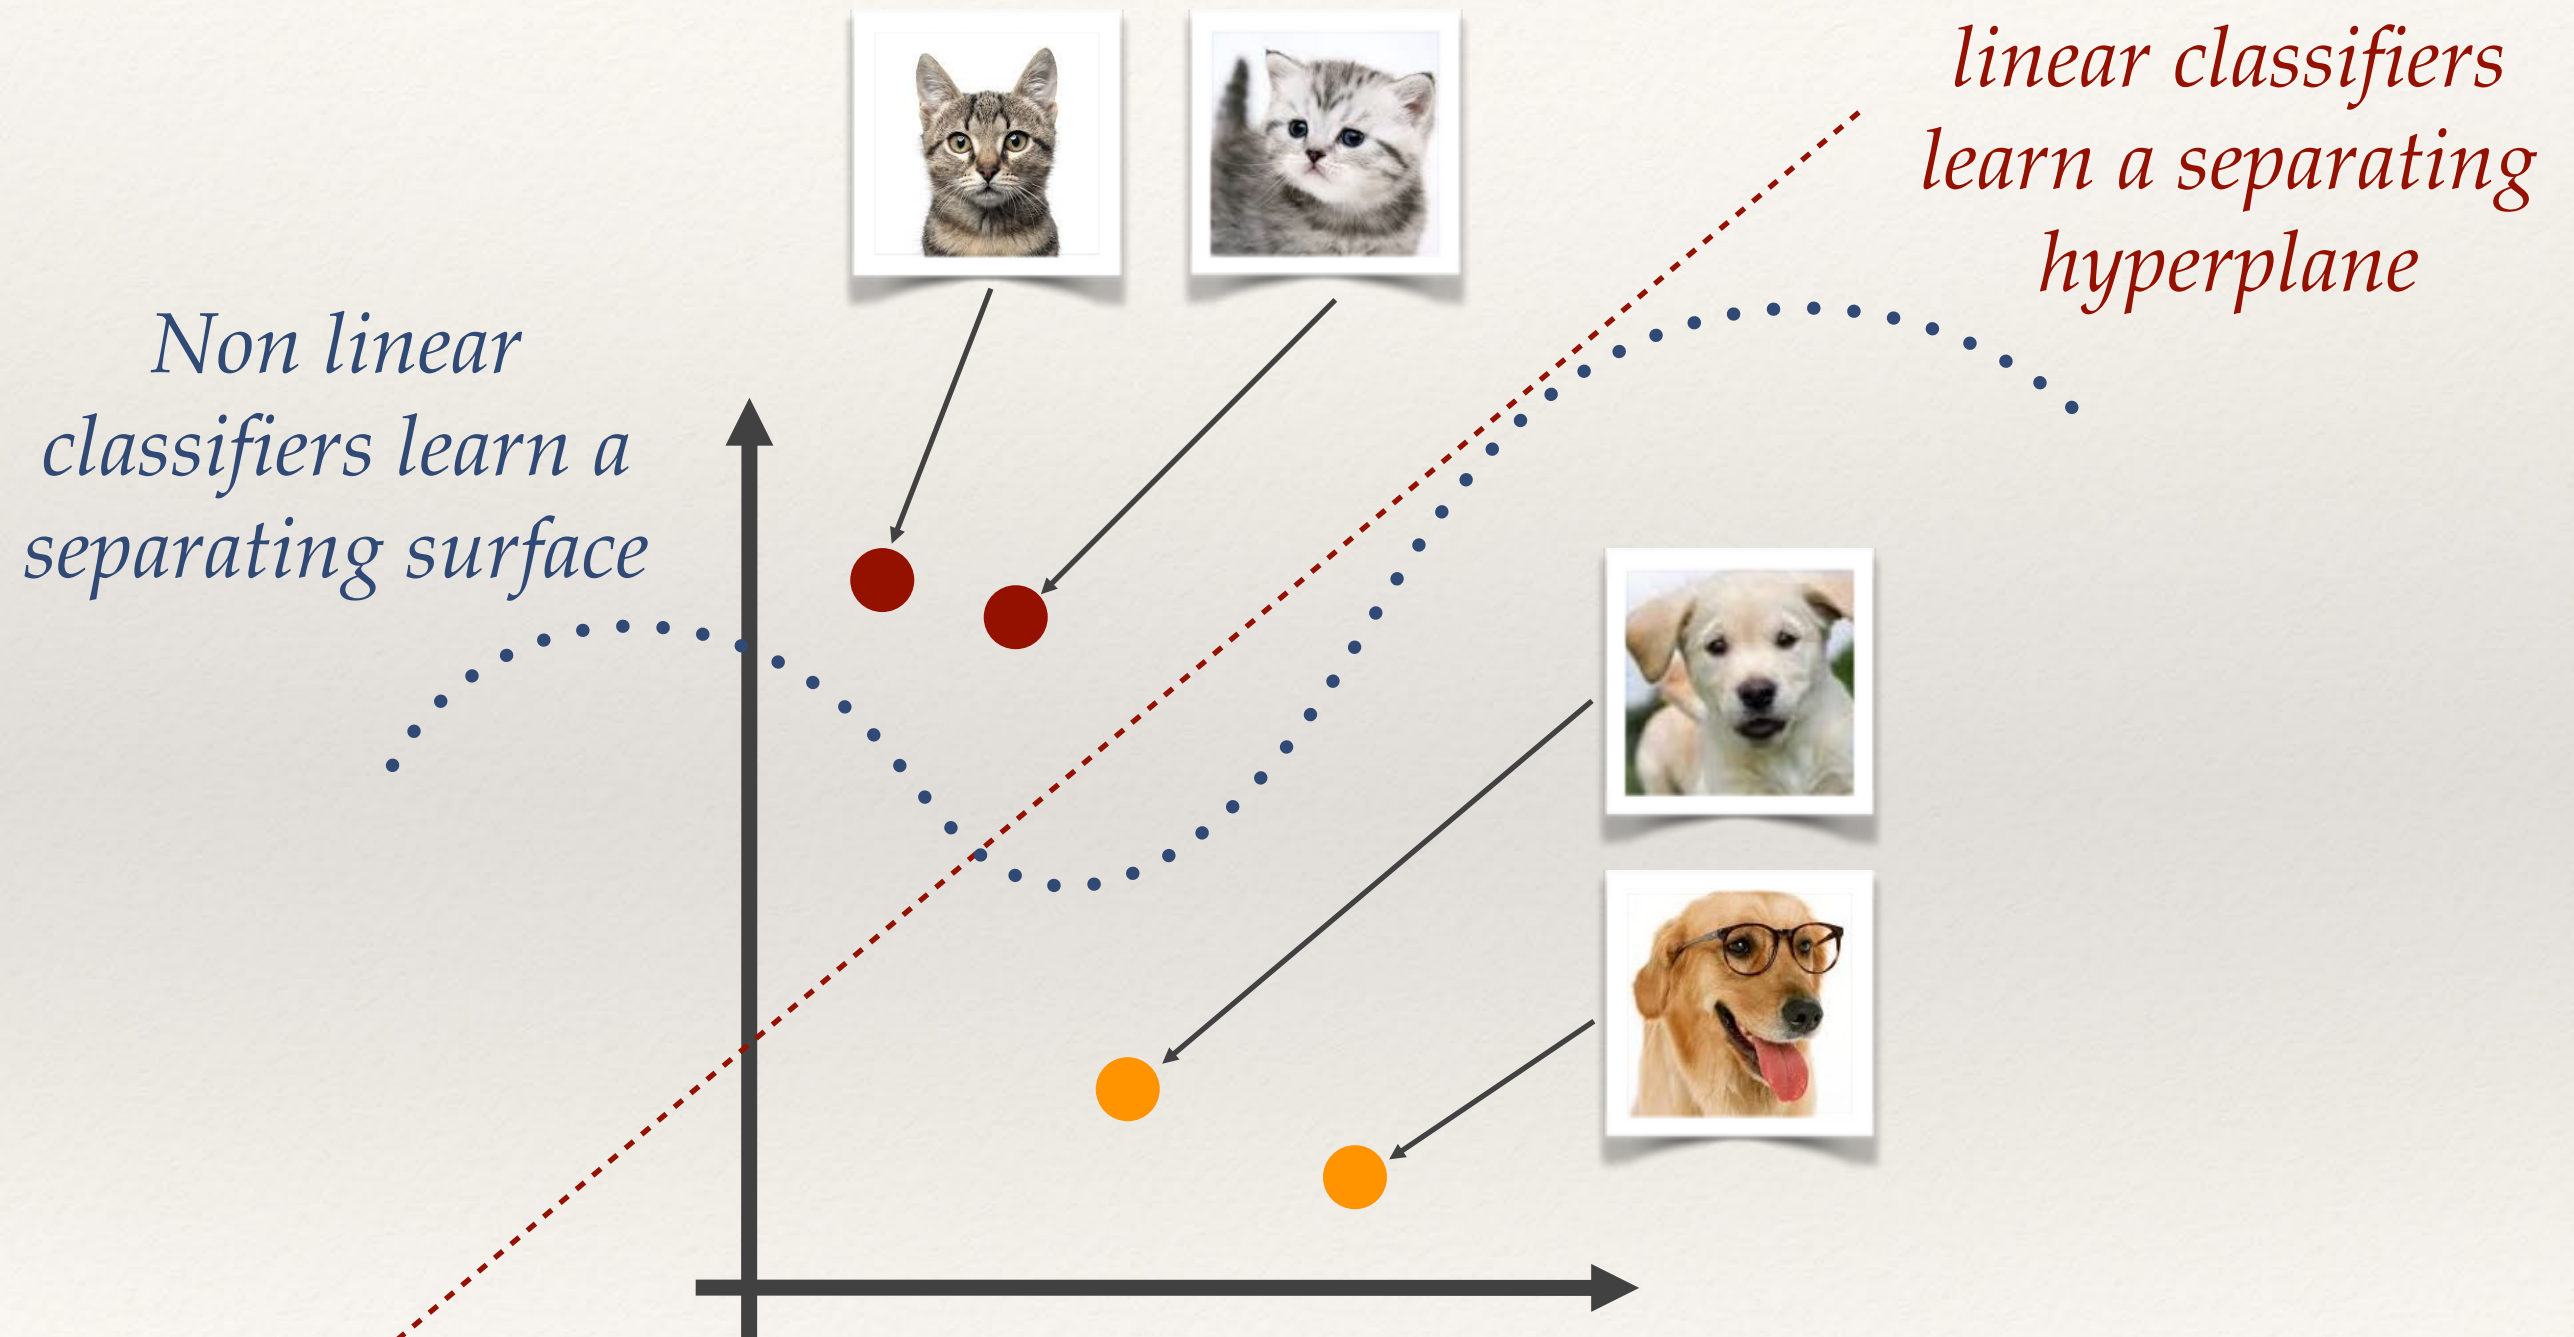
Hansung Kim  
[h.kim@soton.ac.uk](mailto:h.kim@soton.ac.uk)



# Recap: Computer Vision Systems

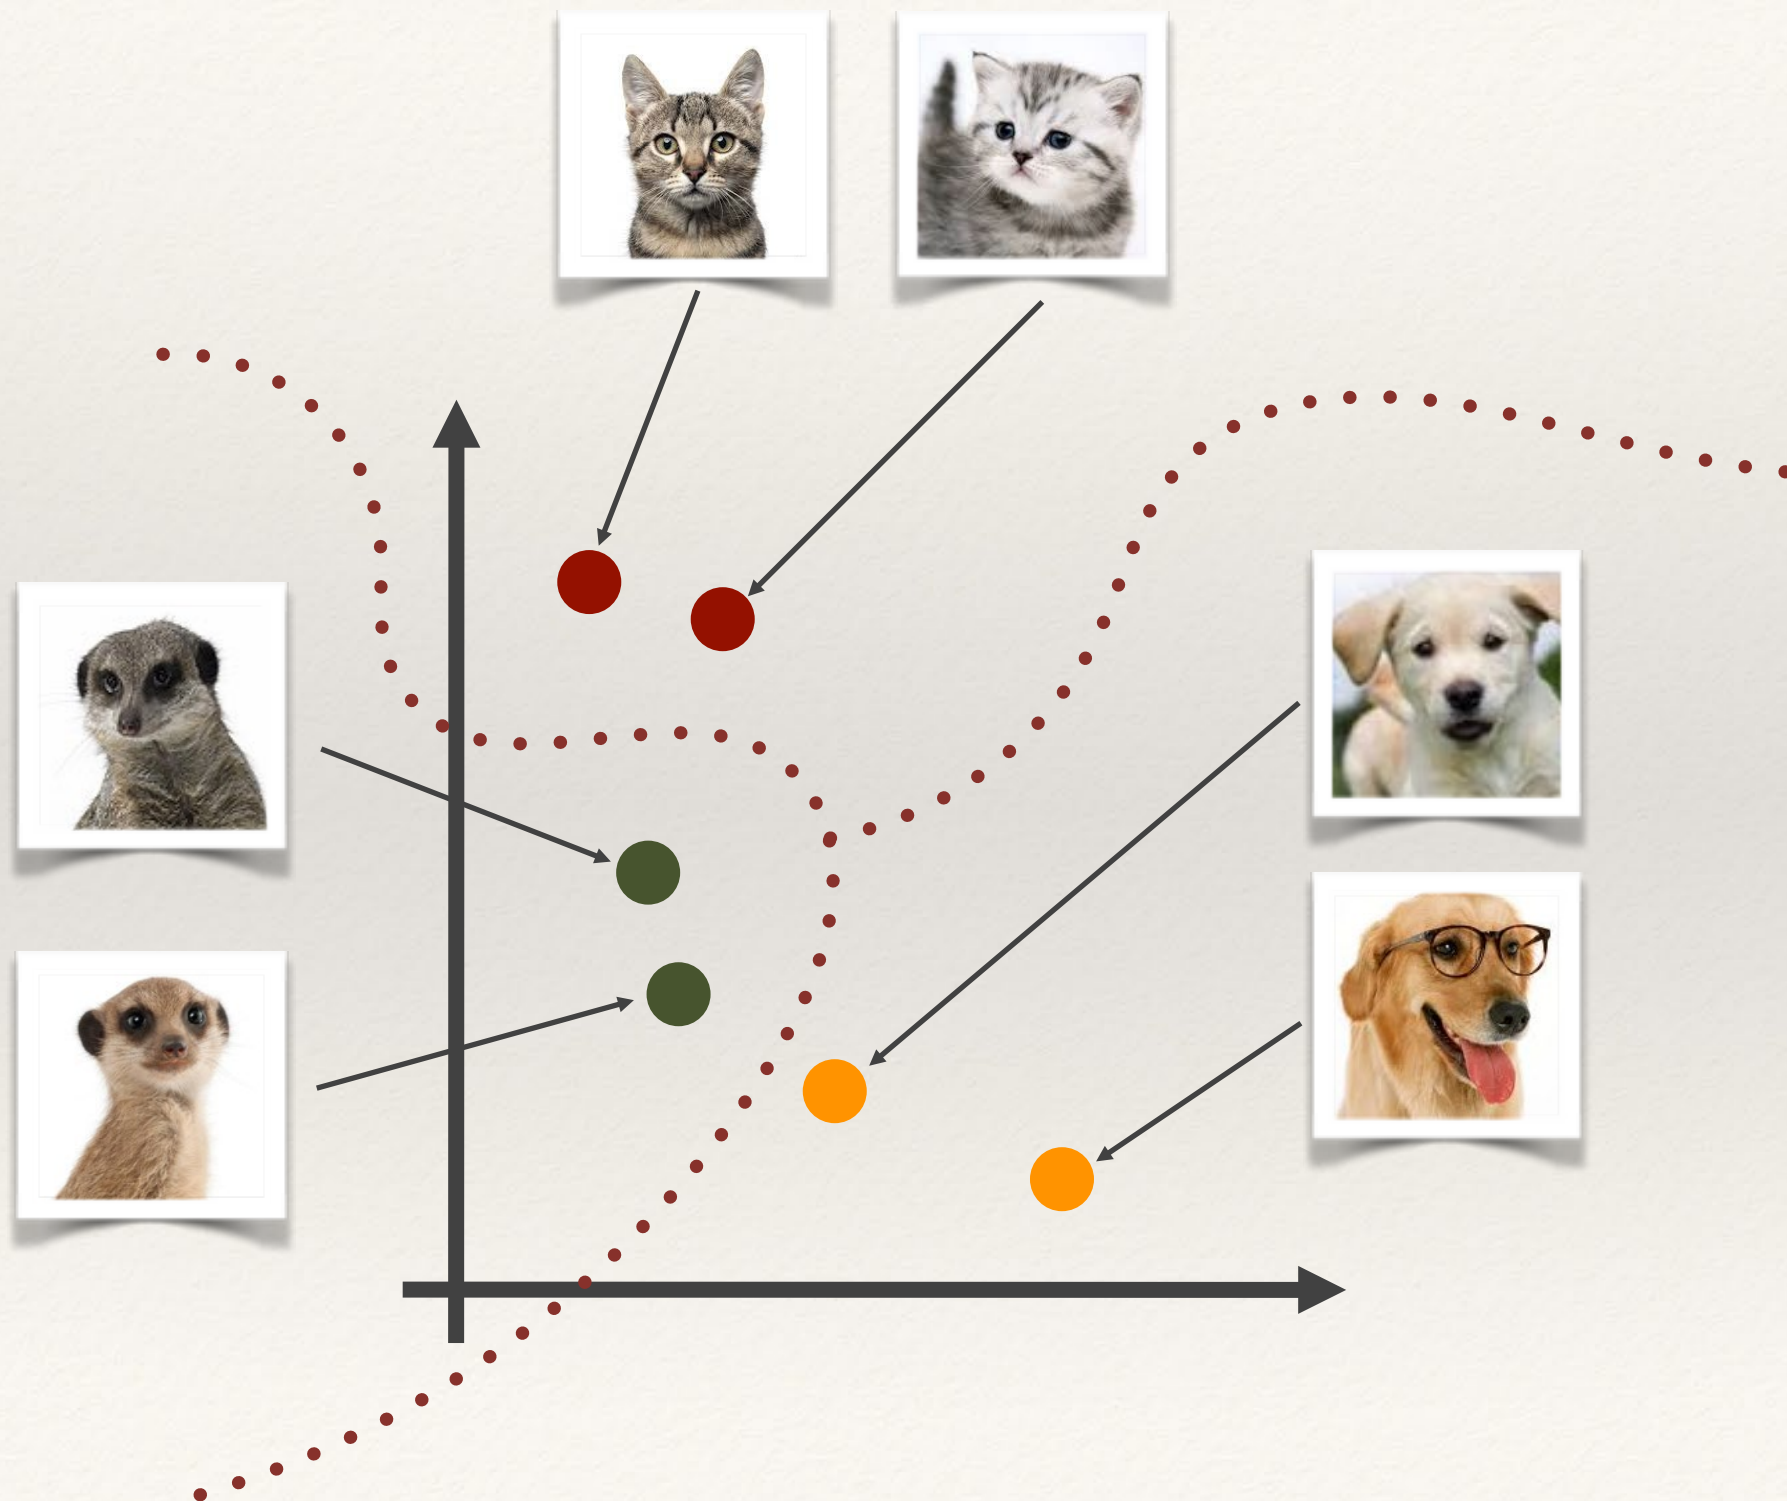


# Recap: Binary classification





# Recap: Multi-class classification



# Multilabel classification

CAT

DOG



in the context of images often called Automatic Annotation



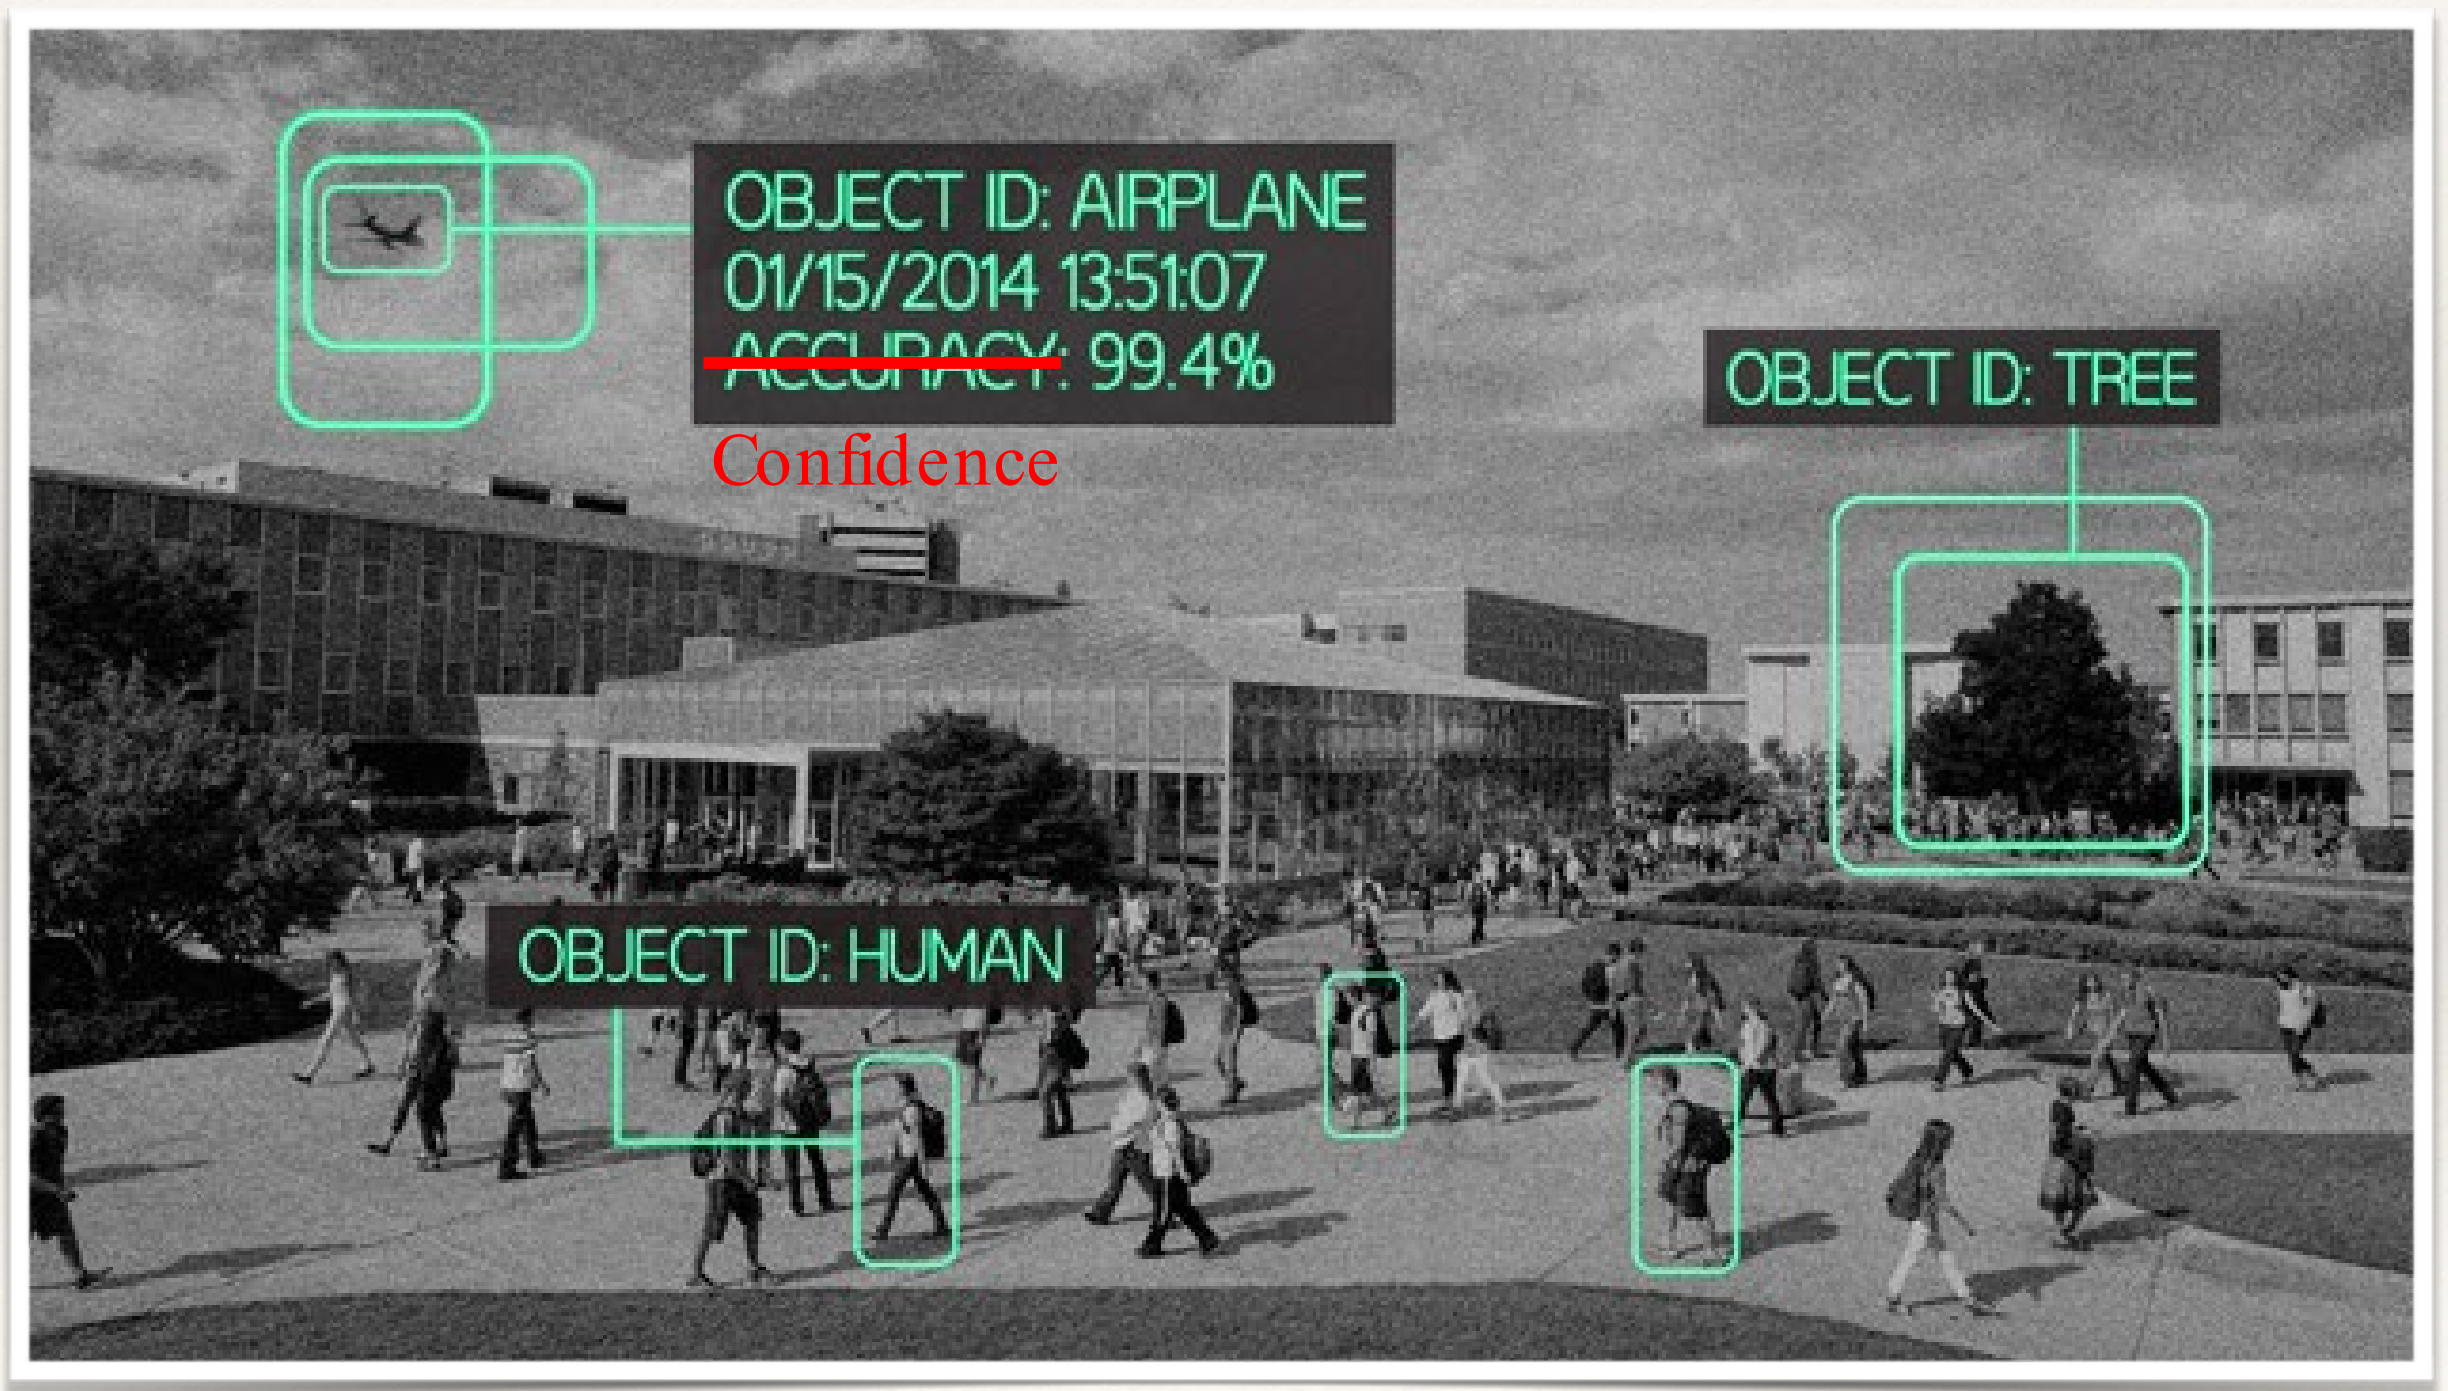


# Object Detection/Localisation



# Challenges in Computer Vision

# Object Recognition in natural scenes





# Scene/Activity Classification

Interacting with a computer



Photographing



Playing music



Riding bike



Riding horse



Running



Walking





# Automatic Annotation (2008)



*sun, bay, sunset, sea, carpet*



*perch, moose, mist, column, ruins*



*pillar, shadows, floor, sea, writing*



*flowers, garden, insect, tulip, blossoms*



*angelfish, mushrooms, fish, coral, fan*



*reptile, sidewalk, pole, detail, hawaii*



*jeep, pair, face, pepper, model*



*tomb, figures, castle, courtyard, fawn*



*detail, pole, church, fountain, window*



*pool, swimmers, athlete, butterfly, people*



*remains, penguin, seals, iguana, marine*



*garden, house, pond, bench, landscape*



# Automatic Annotation with Deep Learning






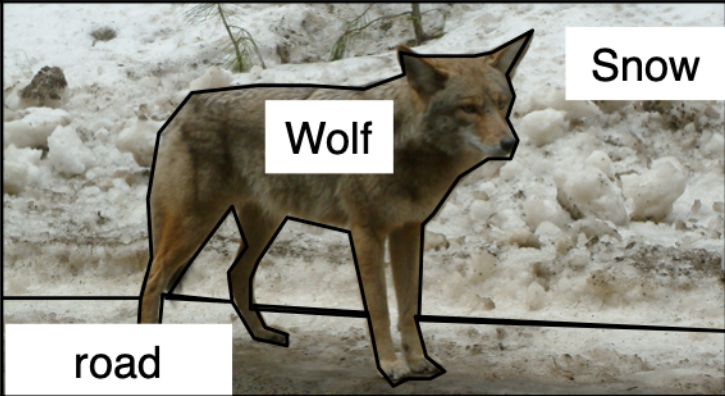
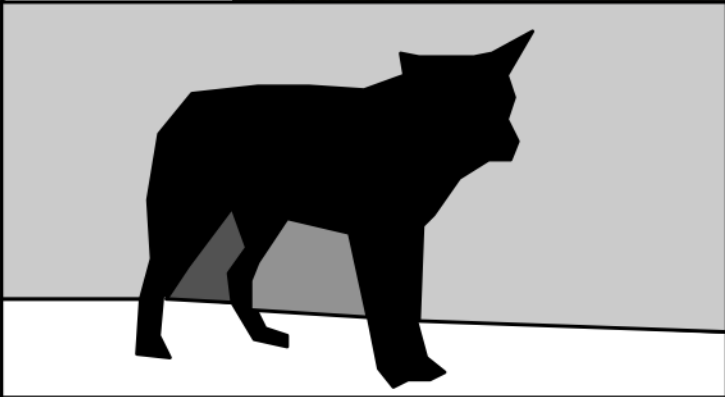

# Real-time Human Detection and Tracking with Deep Learning





# The fundamental problem of computer vision: **The Semantic Gap**



<b>Semantics</b> object relationships and more	Wolf on Road with Snow on Roadside in Yosemite National Park, California on 24/1/2004 at 23:19:11GMT
<b>Object Labels</b> symbolic names of objects	
<b>Objects</b> prototypical combinations of descriptors	
<b>Descriptors</b> feature-vectors	Segmented blobs, Salient regions, Pixel-level histograms, Fourier descriptors, etc...
<b>Raw Media</b> images	





**A car parked on double yellow lines**

# A potted history



---

# Object Recognition

---

- ❖ 1999 - SIFT matching
  - ❖ Very powerful, but computationally demanding
- ❖ 2001 - Cascades of Haar-like features
  - ❖ Very popular for face detection
- ❖ 2006 - SURF matching
  - ❖ Combined ideas from SIFT and the integral images used for computing Haar-like features

Interest in auto-annotation grew from the late 90s

Bags of “Visual Words” were rather important!

(but not in the same way)



---

# Aside: Optimal codebook size


---

- ❖ The codebook vocabulary needs to be much smaller than for doing image search
  - ❖ In general, machine-learning techniques need much smaller vectors (for both performance and effectiveness)
  - ❖ The visual words can be allowed to be less distinctive, allowing a little more variation between matching features.
  - ❖ Typically, the number of visual words might be as small as a few hundred, and up to a few thousand.





# Machine Translation (2002)



Visual words!



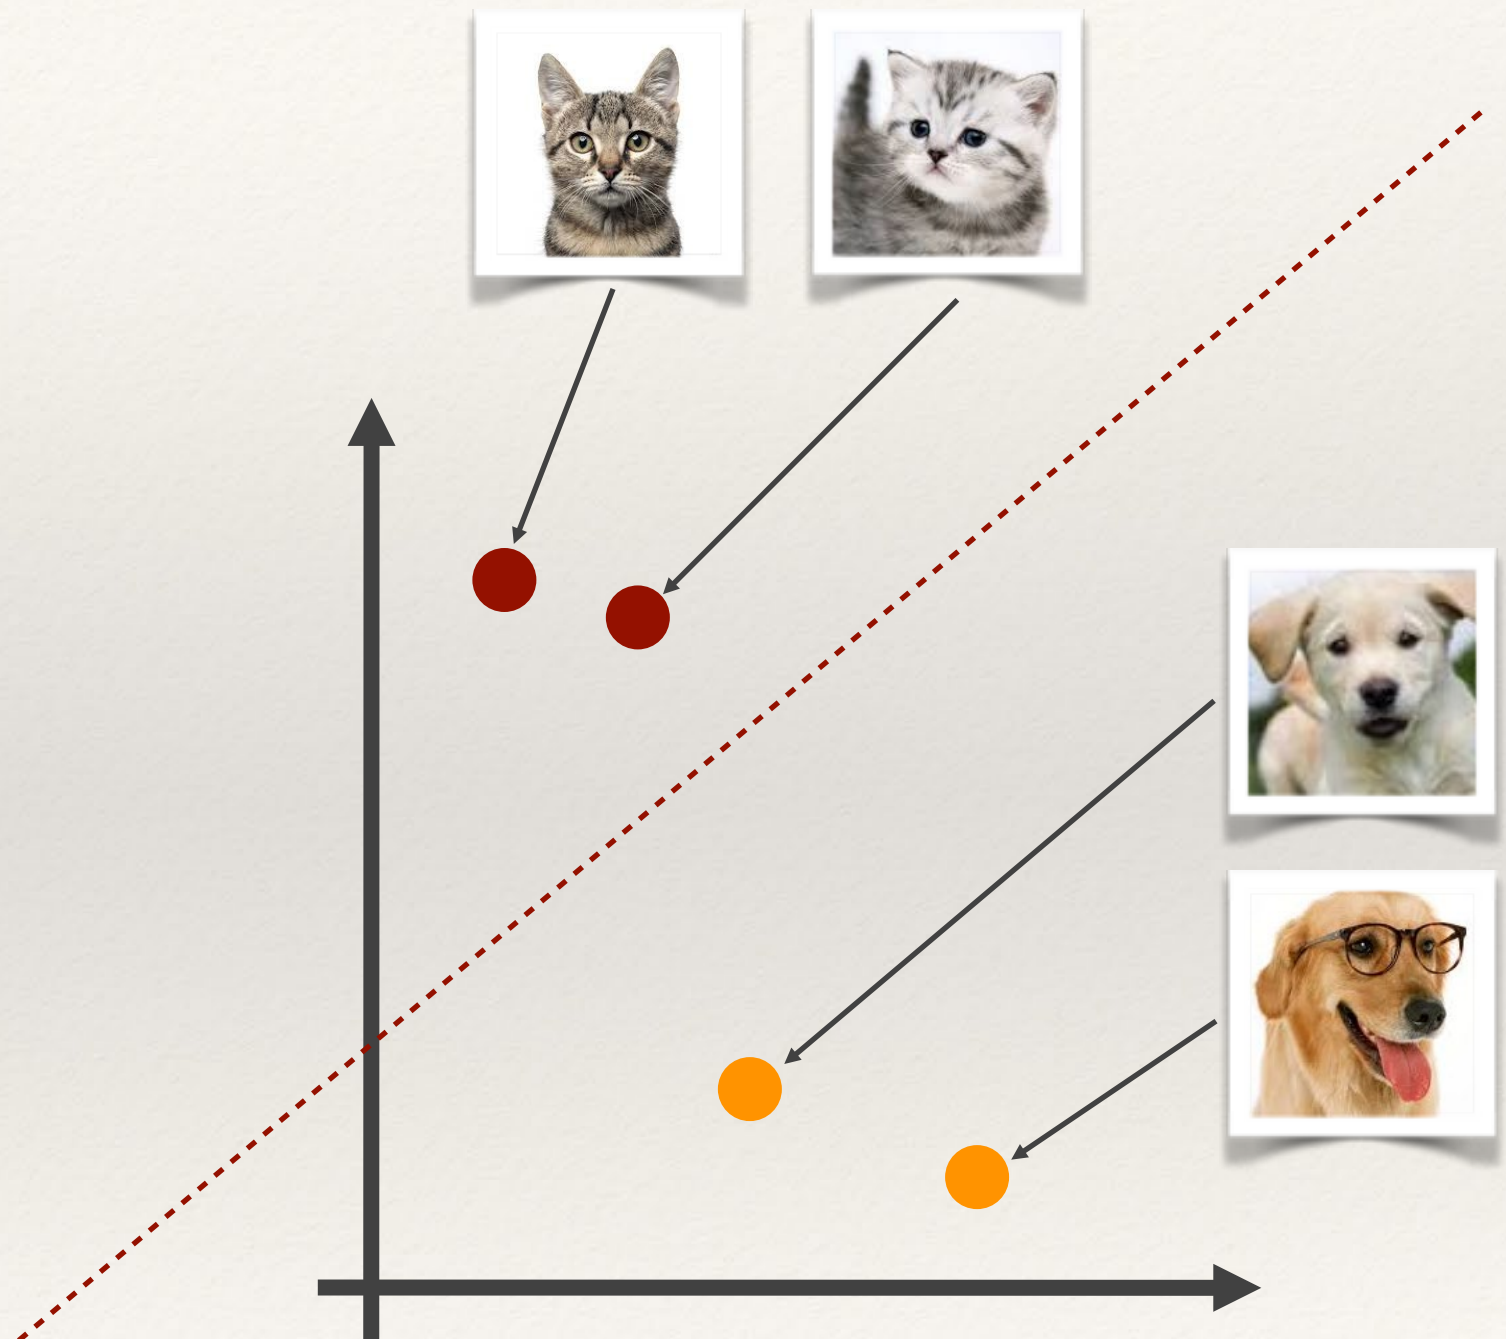
Research focus shifted a little to use of **bigger datasets** in the mid-late 2000s.

Interest in simpler (but more scalable) classifiers grew



# Classifying with BoVW

- ❖ BoVW histogram representations are incredibly useful for image classification and object detection
- ❖ Commonly used with fast linear classifiers and SVMs




Over time the features used to create BoVW representations have improved



# Early global colour visual terms

- ❖ Consider each pixel as a visual word based on the quantisation of its colour to a discrete set of values.
- ❖ The BoVW Histogram is just a joint colour histogram that we saw earlier



# Visual words from regions/segments



[ 1 2 0 0 6 ]




# Visual words from interest points

## Salient region detection



## Local Descriptors

0,255,1,...  
40,1,188,...  
122,32,44,...  
54,231,123...  
121,240,199,...  
123,241,190,...



**Vocabulary of visual terms learnt through hierarchical k-means**

**Vector Quantisation**

## Word Occurrence Vectors


im1: 0,1,2,0,0,1,1,0,1  
im2: 0,1,0,0,1,0,1,0,1  
im3: 2,0,1,1,0,1,0,2,0  
...

Local features extracted around interest points work okay for classification, but there are more recent strategies that can work better...

***Densely sampled features***

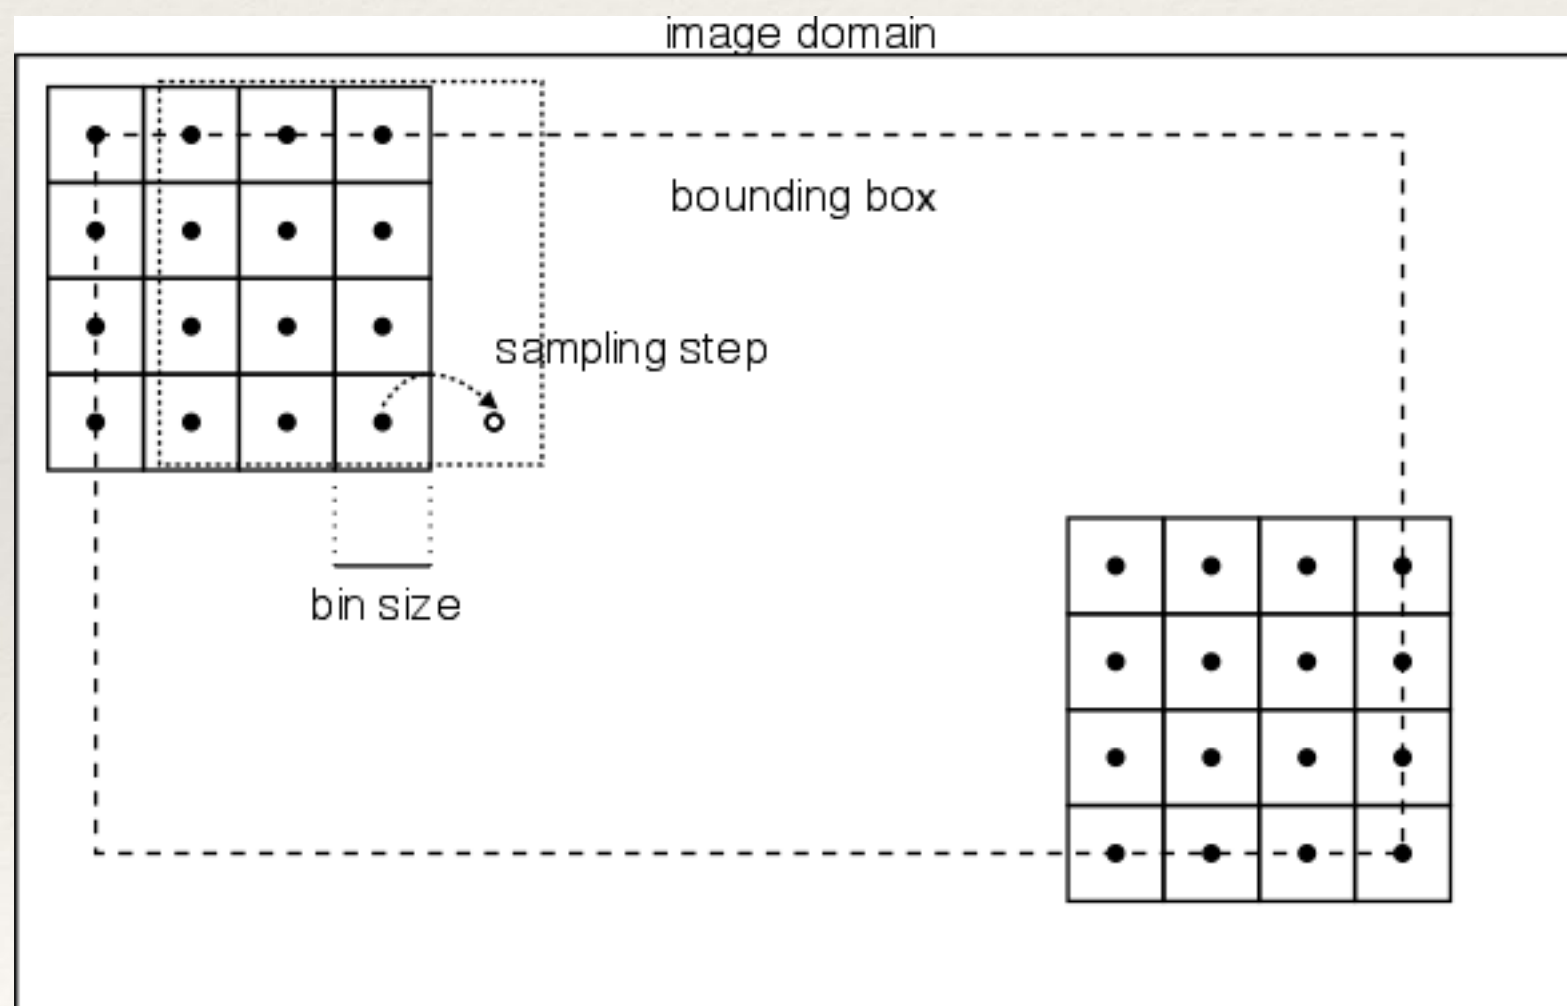


# Dense Local Image Patches



# Dense SIFT


*Rather than extracting your SIFT features at DoG interest points, you could extract them across a dense grid - this gives much more coverage of the entire image.*



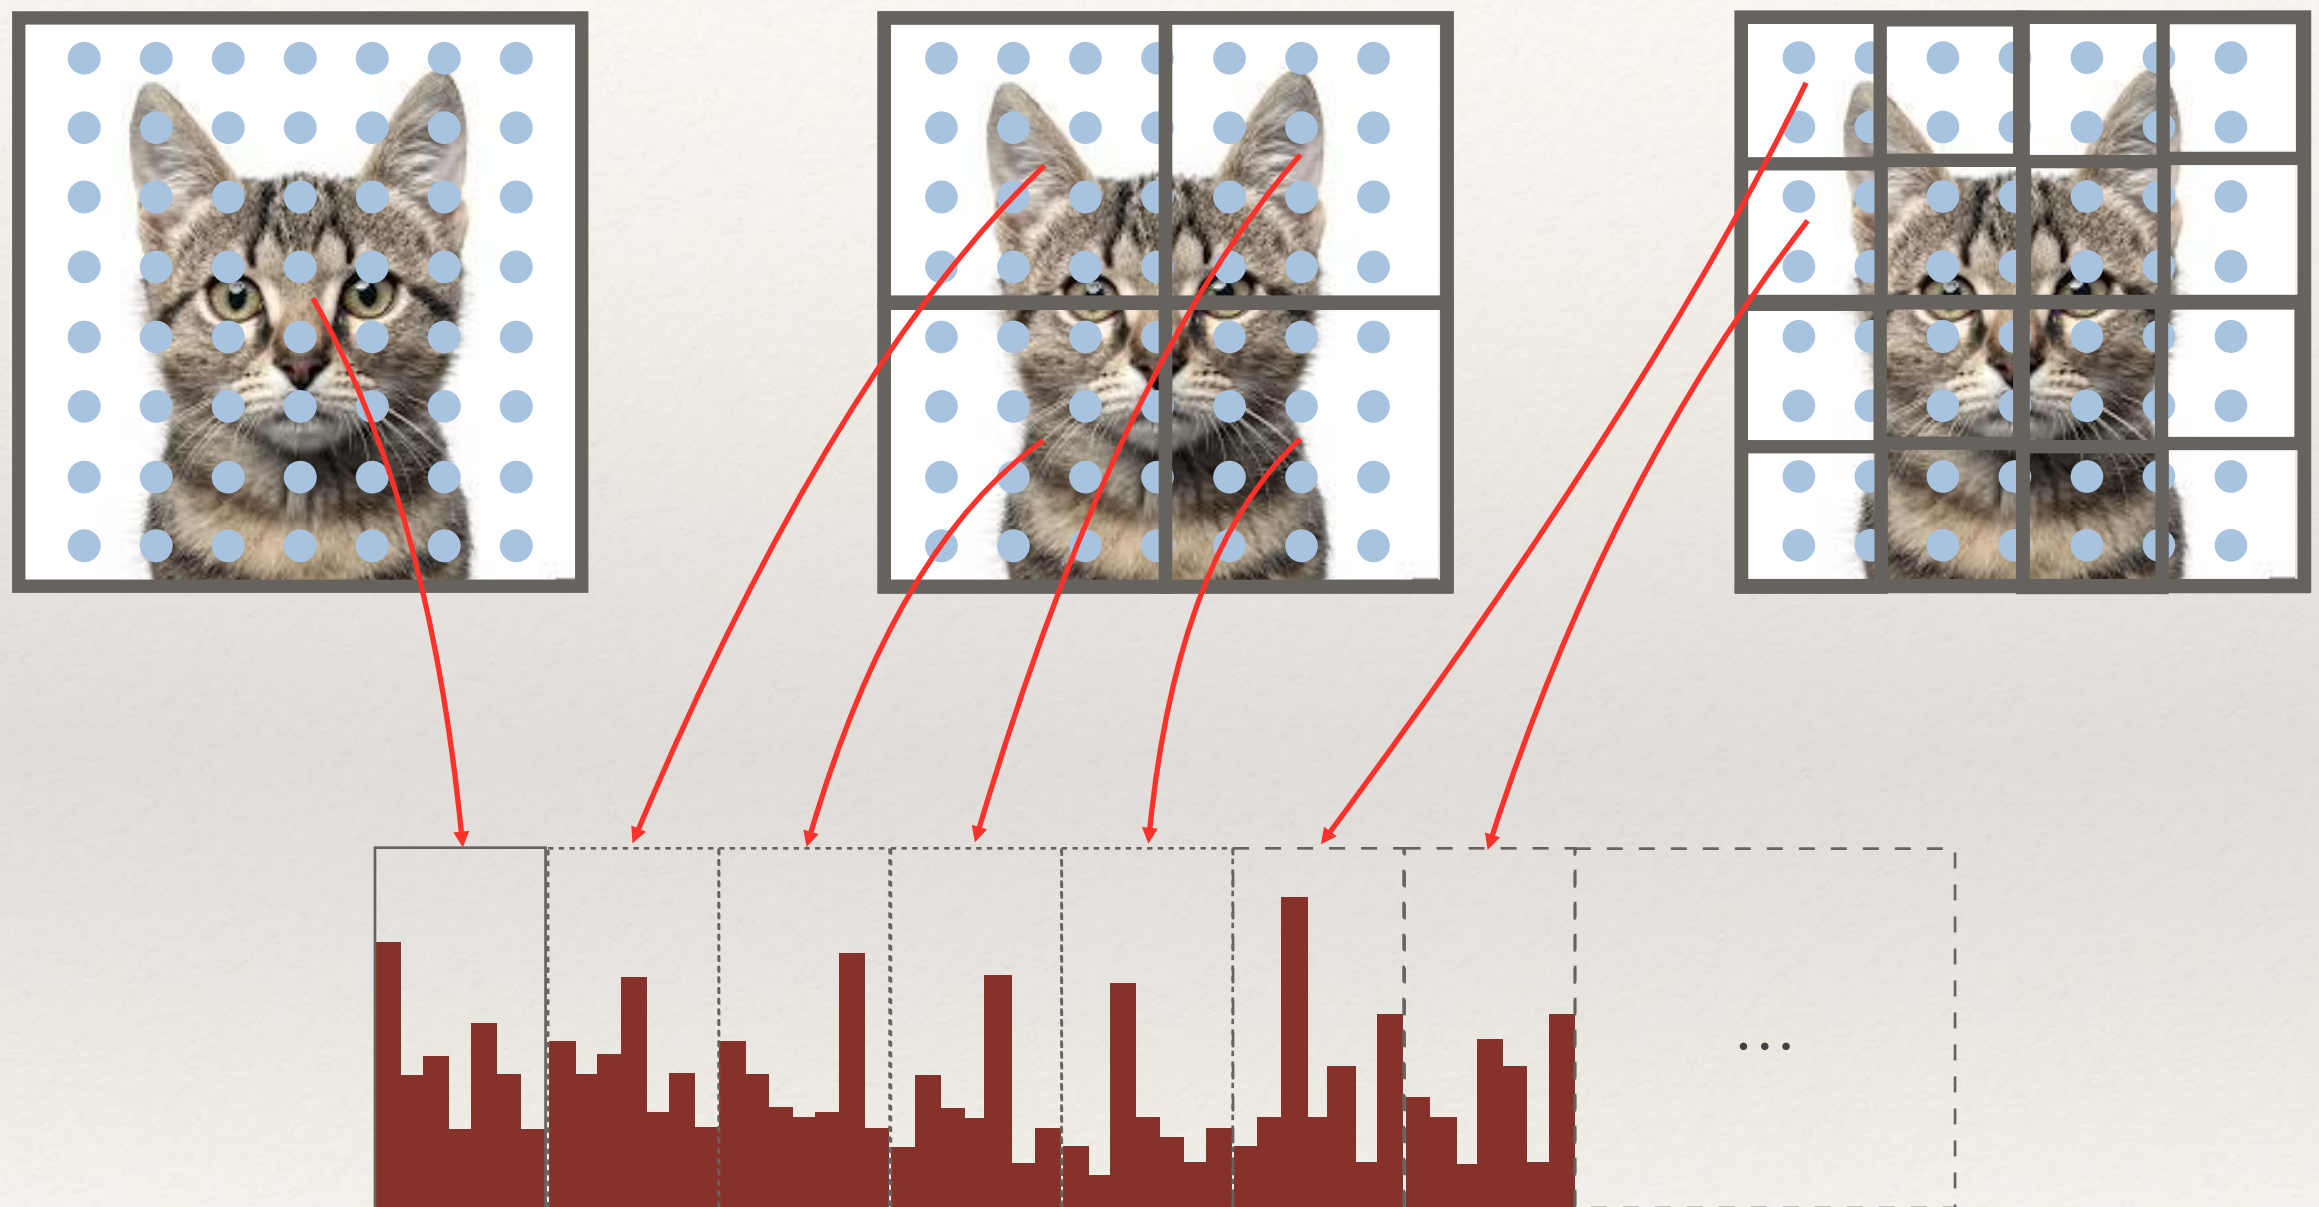


# Pyramid Dense SIFT

- ❖ For even better performance and coverage, you can sample in a Gaussian pyramid
- ❖ Note that the sampling region is a fixed size, so at higher scales you sample more content



# Spatial Pyramids



*PHOW: Pyramid Histogram of Words = Hist(VQ(Pyramid Dense SIFT)) + Spatial Pyramid*



# Developing and benchmarking a BoVW scene classifier

---

# Evaluation Dataset

---

- ❖ Common for academic research to use standardised datasets for developing scene classifiers and comparing results
- ❖ Datasets are usually split into labelled “training” and “test” sets.
  - ❖ Only the training set can be used to train the classifier
  - ❖ Sometimes the test set labels are *withheld* completely to ensure there is no cheating!





---

# Building the BoVW

---

- ❖ Firstly the raw features need to be extracted from the training images
- ❖ Then (if necessary) learn a codebook from these features
  - ❖ i.e. using k-means on the raw features
    - ❖ might be a *uniform random sample* of all the features rather than all of them
- ❖ Apply (vector) quantisation to the raw features and count the number of occurrences to build histograms for each image



---

# Training classifiers

---

- ❖ Classifiers can be trained using the histograms.
  - ❖ e.g. OvR linear classifiers
- ❖ You might train on a subset of the training data (cross-validation)
  - ❖ and use the remaining data to “validate” and optimise parameters.
  - ❖ Once you’ve chosen the optimal parameters you can then re-train using the optimal values.





---

# Classifying the test set

---

- ❖ You're now in a position to apply the classifiers to the test data:
  - ❖ Extract the features
  - ❖ Quantise the features (using the codebook developed from the training set!)
  - ❖ Compute the occurrence histograms
  - ❖ Use the classifiers to find the most likely class



---

# Evaluating Performance

---

- ❖ Lots of ways to evaluate performance of classification on the test (and validation) set.
- ❖ Conceptually the simplest summary measure is probably *average precision*
- ❖ this is literally the proportion of number of correct classifications to the total number of predictions





---

# Summary

---

- ❖ Object recognition, scene classification and automatic annotation are all important tasks in computer vision.
- ❖ Researchers are striving to narrow the “semantic gap” between what computers can perceive compare to humans.
- ❖ The BoVW approach lends itself to high-performance image classification
  - ❖ Performance is increased if the local features are sampled densely

# The Final Coursework



---

# Further reading

---

- ❖ Wikipedia has good articles
  - ❖ Vector quantisation:  
[http://en.wikipedia.org/wiki/Vector\\_quantization](http://en.wikipedia.org/wiki/Vector_quantization)
  - ❖ Bag of Visual Words (and applications):  
[http://en.wikipedia.org/wiki/Bag-of-words\\_model\\_in\\_computer\\_vision](http://en.wikipedia.org/wiki/Bag-of-words_model_in_computer_vision)
- ❖ First work on spatial pyramids:
  - ❖ [http://ieeexplore.ieee.org/xpls/abs\\_all.jsp?arnumber=1641019&tag=1](http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=1641019&tag=1)
- ❖ Info on the homogeneous kernel map (including software implementations and papers):  
<http://www.robots.ox.ac.uk/~vgg/software/homkernmap/>
- ❖ Practical exercises
  - ❖ Chapter 12 of the OpenIMAJ tutorial covers dense local feature extraction, spatial pyramids and fast linear classification for learning a set of 101 object categories (Not only for Java).