
Algorithmic Game Theory

COMP6207

Lecture 13: Hospitals/Residents Problem

Baharak Rastegari

b.rastegari@soton.ac.uk

Electronics and Computer Science
University of Southampton

Learning Outcomes

- By the end of this session, you should be able to
 - **Describe** the hospital/residents (HR) problems and its variants
 - **Describe** the relationship between hospital-optimal matchings and resident-optimal matchings
 - **Identify** blocking pairs in instances of HR
 - **Compute** the matching produced by Resident-oriented GS (RGS)
 - **Compute** the matching produced by Hospital-oriented GS (HGS)
 - **Describe** the implications of couples participating in HR

Extensions of Stable Matching problem

- Agents may declare some candidates unacceptable
→ **S**table **M**atching problem with **I**ncomplete lists (**SMI**)
- Agents may be indifferent among several candidates
→ **S**table **M**atching problem with **T**ies (**SMT**)
- Both incomplete lists and indifferences are allowed
→ **S**table **M**atching problem with **T**ies and **I**ncomplete lists (**SMTI**)

Today

- Agents on one side can get matched to several candidates
 - Many-one stable matching problem
 - **H**ospitals/**R**esidents problem (**HR**) and HR with **T**ies (**HRT**)

Note: In **HR** incomplete lists are allowed.

From Stable Marriage to the **Hospitals/Residents** problem



Match Day 2017. Credit: Charles E. Schmidt College of Medicine, FAU.
For more photos of this important day of medical students' life click [here](#).

and its variants

How it works in practice, usually

- Junior doctors (or residents) must undergo training in hospitals
- Applicants rank hospitals in order of preference
- Hospitals do likewise with their applicants
- Centralised matching schemes (clearinghouses) produce a matching in several countries
 - US (National Resident Matching Program)
 - Canada (Canadian Resident Matching Service)
 - Japan (Japan Residency Matching Program)
 - UK (UK Foundation Programme Office)
- Stability is the key property of a matching
 - [Roth, 1984]

Many-One variants of SMI and SMTI

- Agents on one side can get matched to several candidates
 - Many-one stable matching problem
 - **H**ospitals/**R**esidents problem (**HR**) and HR with **T**ies (**HRT**)

Note: In **HR** incomplete lists are allowed.

Hospitals / Residents problem (HR)

- **Participants**

- A set of n_1 residents $\{r_1, r_2, \dots, r_{n_1}\}$
- A set of n_2 hospitals $\{h_1, h_2, \dots, h_{n_2}\}$

- Each hospital has a **capacity**

- **Preferences**

- Residents rank **acceptable** hospitals in strict order of preference, hospitals do likewise
- We assume that unacceptability is **mutual**: if **h** finds **r** unacceptable then **r** finds **h** unacceptable and vice versa

Matching in HR

A *matching* M is a set of resident-hospital pairs such that:

- $(r, h) \in M \Rightarrow r, h$ find each other acceptable
- No resident appears in more than one pair
- No hospital appears in more pairs than its capacity

HR: example matching

$\mathbf{r}_1: \mathbf{h}_2 \mathbf{h}_1$

$\mathbf{r}_2: \mathbf{h}_1 \mathbf{h}_2$

$\mathbf{r}_3: \mathbf{h}_1 \mathbf{h}_3$

$\mathbf{r}_4: \mathbf{h}_2 \mathbf{h}_3$

$\mathbf{r}_5: \mathbf{h}_2 \mathbf{h}_1$

$\mathbf{r}_6: \mathbf{h}_1 \mathbf{h}_2$

Resident preferences

Each hospital has capacity **2**

$\mathbf{h}_1: \mathbf{r}_1 \mathbf{r}_3 \mathbf{r}_2 \mathbf{r}_5 \mathbf{r}_6$

$\mathbf{h}_2: \mathbf{r}_2 \mathbf{r}_6 \mathbf{r}_1 \mathbf{r}_4 \mathbf{r}_5$

$\mathbf{h}_3: \mathbf{r}_4 \mathbf{r}_3$

Hospital preferences

HR: example matching

r_1 : h_2 h_1

r_2 : h_1 h_2

r_3 : h_1 h_3

r_4 : h_2 h_3

r_5 : h_2 h_1

r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**

h_1 : r_1 r_3 r_2 r_5 r_6

h_2 : r_2 r_6 r_1 r_4 r_5

h_3 : r_4 r_3

Hospital preferences

$M = \{(r_1, h_1), (r_2, h_2), (r_3, h_3), (r_5, h_2), (r_6, h_1)\}$ (size **5**)

HR: stability

- Matching M is *stable* if M admits no *blocking pair*
 - (r, h) is a blocking pair of matching M if:
 1. r, h find each other acceptable
and
 2. *either* r is unmatched in M
or r prefers h to his/her assigned hospital in M
and
 3. *either* h is undersubscribed in M
or h prefers r to its worst resident assigned in M

HR: blocking pair (1)

$r_1: h_2 \text{ } \textcircled{h_1}$
 $r_2: \text{ } \textcircled{h_1} \text{ } \textcircled{h_2}$
 $r_3: h_1 \text{ } \textcircled{h_3}$
 $r_4: h_2 \text{ } h_3$
 $r_5: \textcircled{h_2} \text{ } h_1$
 $r_6: \textcircled{h_1} \text{ } h_2$

Resident preferences

Each hospital has capacity **2**

$h_1: \textcircled{r_1} r_3 \text{ } \textcircled{r_2} r_5 \textcircled{r_6}$
 $h_2: \textcircled{r_2} r_6 r_1 r_4 \textcircled{r_5}$
 $h_3: r_4 \textcircled{r_3}$

Hospital preferences

$M = \{(r_1, h_1), (r_2, h_2), (r_3, h_3), (r_5, h_2), (r_6, h_1)\}$ (size **5**)

(r_2, h_1) is a blocking pair of M

HR: blocking pair (2)

$r_1: h_2$ (h_1)
 $r_2: h_1$ (h_2)
 $r_3: h_1$ (h_3)
 $r_4: (h_2)$ h_3
 $r_5: (h_2)$ h_1
 $r_6: (h_1)$ h_2

Resident preferences

Each hospital has capacity **2**

$h_1: (r_1) r_3 r_2 r_5 (r_6)$
 $h_2: (r_2) r_6 r_1 (r_4) (r_5)$
 $h_3: r_4 (r_3)$

Hospital preferences

$M = \{(r_1, h_1), (r_2, h_2), (r_3, h_3), (r_5, h_2), (r_6, h_1)\}$ (size **5**)

(r_4, h_2) is a blocking pair of M

HR: blocking pair (3)

$r_1: h_2$ (h_1)
 $r_2: h_1$ (h_2)
 $r_3: h_1$ (h_3)
 $r_4: h_2$ (h_3)
 $r_5: (h_2)$ h_1
 $r_6: (h_1)$ h_2

Resident preferences

Each hospital has capacity **2**

$h_1: (r_1) r_3 r_2 r_5 (r_6)$
 $h_2: (r_2) r_6 r_1 r_4 (r_5)$
 $h_3: (r_4) (r_3)$

Hospital preferences

$M = \{(r_1, h_1), (r_2, h_2), (r_3, h_3), (r_5, h_2), (r_6, h_1)\}$ (size **5**)

(r_4, h_3) is a blocking pair of M

HR: stable matching

r_1 : h_2 h_1
 r_2 : h_1 h_2
 r_3 : h_1 h_3
 r_4 : h_2 h_3
 r_5 : h_2 h_1
 r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**

h_1 : r_1 r_3 r_2 r_5 r_6
 h_2 : r_2 r_6 r_1 r_4 r_5
 h_3 : r_4 r_3

Hospital preferences

$M = \{(r_1, h_1), (r_2, h_2), (r_3, h_3), (r_5, h_2), (r_6, h_1)\}$ (size **5**)

r_5 is unmatched
 h_3 is undersubscribed

Stable Matching in HR

A *matching* M is a set of resident-hospital pairs such that:

- $(r, h) \in M \Rightarrow r, h$ find each other acceptable
- No resident appears in more than one pair
- No hospital appears in more pairs than its capacity

Matching M is *stable* if M admits no *blocking pair*

– (r, h) is a blocking pair of matching M if:

1. r, h find each other acceptable

and

2. *either* r is unmatched in M
or r prefers h to his/her assigned hospital in M

and

3. *either* h is undersubscribed in M
or h prefers r to its worst resident assigned in M

Resident-oriented GS (RGS)

An extension of Gale-Shapley to HR

Resident-oriented GS (residents, hospitals, capacities, preferences)

```
1  M =  $\emptyset$ ;  //assign all residents and hospitals to be free
2  While (some resident  $r_i$  is unmatched and has a non-empty list)
3       $r_i$  applies to the first hospital  $h_j$  on her list;
4      M = M  $\cup$  {( $r_i$ ,  $h_j$ )};
5      If ( $h_j$  is over-subscribed)
6           $r_k$  = worst resident assigned to  $h_j$ ;
7          M = M  $\setminus$  {( $r_k$ ,  $h_j$ )}; //  $r_k$  is set free
8      If ( $h_j$  is full)
9           $r_k$  = worst resident assigned to  $h_j$ ;
10         For (each successor  $r_1$  of  $r_k$  on  $h_j$ 's list)
11             delete  $r_1$  from  $h_j$ 's list;
12             delete  $h_j$  from  $r_1$ 's list;
13  output the engaged pairs, who form a stable matching;
```

RGS algorithm: example

r_1 : h_2 h_1
 r_2 : h_1 h_2
 r_3 : h_1 h_3
 r_4 : h_2 h_3
 r_5 : h_2 h_1
 r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**

h_1 : r_1 r_3 r_2 r_5 r_6
 h_2 : r_2 r_6 r_1 r_4 r_5
 h_3 : r_4 r_3

Hospital preferences

RGS algorithm: example

r_1 : h_2 h_1
 r_2 : h_1 h_2
 r_3 : h_1 h_3
 r_4 : h_2 h_3
 r_5 : h_2 h_1
 r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**

h_1 : r_1 r_3 r_2 r_5 r_6
 h_2 : r_2 r_6 r_1 r_4 r_5
 h_3 : r_4 r_3

Hospital preferences

Stable matching: $M = \{(r_1, h_2), (r_2, h_1), (r_3, h_1), (r_4, h_3), (r_6, h_2)\}$

Optimal stable matchings

Straightforward extension of results for **SMI**:

Claim. When preferences are **strict**, all executions of resident-oriented GS (RGS) return the **resident-optimal** (stable) matching.

Claim. When preferences are **strict**, all executions of hospital-oriented GS (HGS) return the **hospital-optimal** (stable) matching.

Claim. When preferences are strict, the hospital-optimal matching is the **resident-pessimal** matching, and vice versa.

Hospital-oriented GS (HGS)

An extension of Gale-Shapley to HR

Hospital-oriented GS (residents, hospitals, capacities, preferences)

```
1  M =  $\emptyset$ ; //assign all residents and hospitals to be free
2  While (some hospital  $h_i$  is undersubscribed and has non-empty list)
3       $s_i = q_i - |M(h_i)|$ ; //the number of available seats in  $h_i$ 
4       $h_i$  applies to the first  $s_i$  residents on its list;
5      foreach ( $r_j$  that  $h_i$  applies to)
6          If ( $r_j$  is free)    M = M  $\cup$  {( $r_j$ ,  $h_i$ )};
7          else if ( $r_j$  prefers  $h_i$  to her current hospital  $h'$ )
8              M = M  $\cup$  {( $r_j$ ,  $h_i$ )}  $\setminus$  {( $r_j$ ,  $h'$ )}; //h' is set free
9          If ( $r_j$  is not free)
10             For (each successor  $h_1$  of  $M(r_j)$  on  $r_j$ 's list)
11                 delete  $h_1$  from  $r_j$ 's list;
12                 delete  $r_j$  from  $h_1$ 's list;
13  output the engaged pairs, who form a stable matching;
```

This If statement is not necessary.
We can skip to For loop right away.

HGS algorithm: example

$\mathbf{r}_1: \mathbf{h}_2 \mathbf{h}_1$

$\mathbf{r}_2: \mathbf{h}_1 \mathbf{h}_2$

$\mathbf{r}_3: \mathbf{h}_1 \mathbf{h}_3$

$\mathbf{r}_4: \mathbf{h}_2 \mathbf{h}_3$

$\mathbf{r}_5: \mathbf{h}_2 \mathbf{h}_1$

$\mathbf{r}_6: \mathbf{h}_1 \mathbf{h}_2$

Resident preferences

Each hospital has capacity **2**

$\mathbf{h}_1: \mathbf{r}_1 \mathbf{r}_3 \mathbf{r}_2 \mathbf{r}_5 \mathbf{r}_6$

$\mathbf{h}_2: \mathbf{r}_2 \mathbf{r}_6 \mathbf{r}_1 \mathbf{r}_4 \mathbf{r}_5$

$\mathbf{h}_3: \mathbf{r}_4 \mathbf{r}_3$

Hospital preferences

HGS algorithm: example

r_1 : h_2 h_1
 r_2 : h_1 h_2
 r_3 : h_1 ~~h_3~~
 r_4 : h_2 h_3
 r_5 : h_2 h_1
 r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**

h_1 : r_1 r_3 r_2 r_5 r_6
 h_2 : r_2 r_6 r_1 r_4 r_5
 h_3 : r_4 ~~r_3~~

Hospital preferences

HGS algorithm: example

$r_1: h_2 \text{ } \textcircled{h_1}$
 $r_2: h_1 \text{ } \textcircled{h_2}$
 $r_3: \textcircled{h_1} h_3$
 $r_4: h_2 \text{ } \textcircled{h_3}$
 $r_5: h_2 h_1$
 $r_6: h_1 \text{ } \textcircled{h_2}$

Resident preferences

Each hospital has capacity **2**

$h_1: \textcircled{r_1} \textcircled{r_3} r_2 r_5 r_6$
 $h_2: \textcircled{r_2} \textcircled{r_6} r_1 r_4 r_5$
 $h_3: \textcircled{r_4} r_3$

Hospital preferences

Stable matching: $M = \{(r_1, h_1), (r_2, h_2), (r_3, h_1), (r_4, h_3), (r_6, h_2)\}$

Hospital-optimal matching

RGS algorithm: example

r_1 : h_2 h_1
 r_2 : h_1 h_2
 r_3 : h_1 h_3
 r_4 : h_2 h_3
 r_5 : h_2 h_1
 r_6 : h_1 h_2

Resident preferences

Each hospital has capacity **2**


h_1 : r_1 r_3 r_2 r_5 r_6
 h_2 : r_2 r_6 r_1 r_4 r_5
 h_3 : r_4 r_3

Hospital preferences

Stable matching: $M = \{(r_1, h_2), (r_2, h_1), (r_3, h_1), (r_4, h_3), (r_6, h_2)\}$


Resident-optimal matching

Hospital-optimal vs Resident-optimal



Resident preferences

Each hospital has capacity **2**



Hospital preferences

 Resident-optimal matching

 Hospital-optimal matching

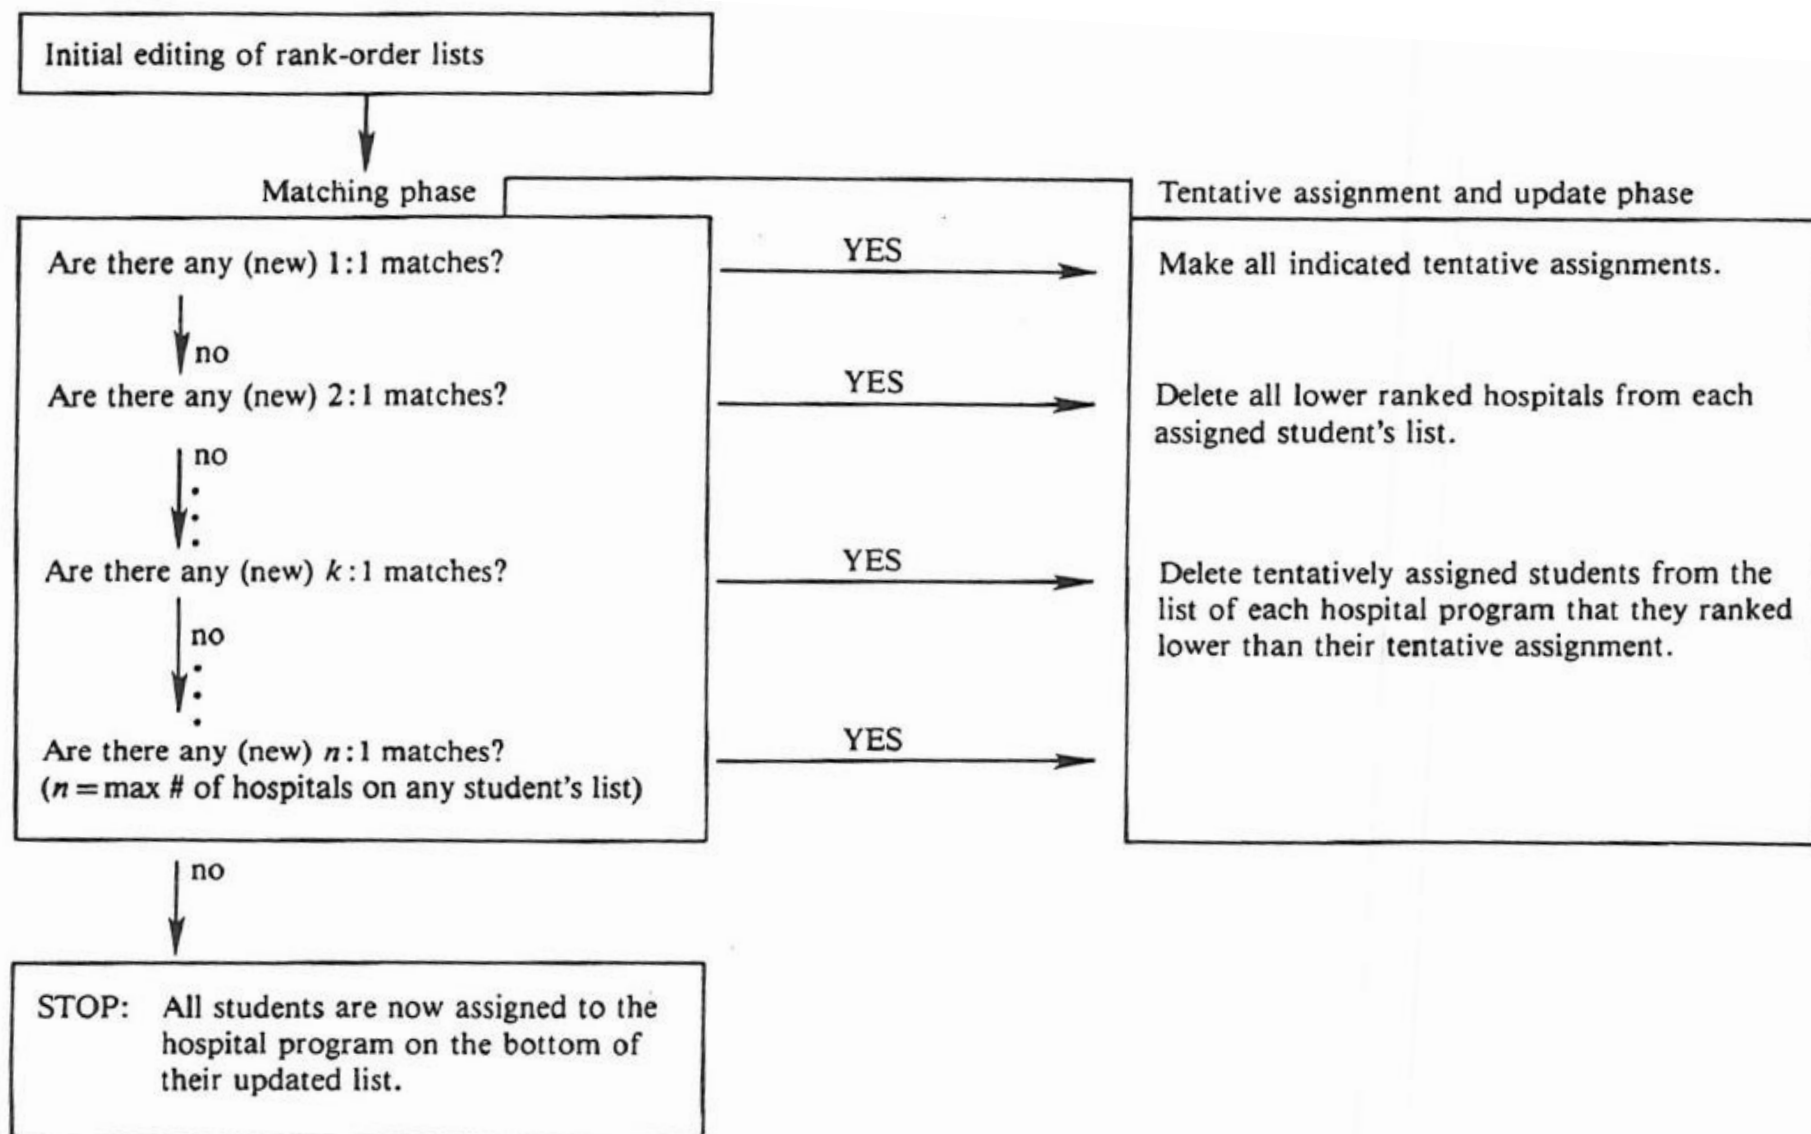


Figure 5.1. The NIMP algorithm.

NIMP (later NRMP) algorithm, adopted in 1951.
Figure from TSM, reproduced from Roth 1984

Rural Hospitals Theorem

Theorem (Roth, 1984; Gale and Sotomayor, 1985; Roth, 1986)

In a given instance of **HR**:

1. the **same residents** are assigned in all stable matchings;
2. each hospital is assigned the **same number of residents** in all stable matchings;
3. any hospital that is **undersubscribed** in one stable matching is assigned **exactly the same set of residents** in all stable matchings.

DS truthfulness

Theorem (Roth; 5.16 in TSM)

A stable matching mechanism that yields the resident-optimal matching makes it a dominant strategy for all residents to state their true preferences.

- Under strict preferences, RGS is dominant-strategy truthful for residents.

Theorem (Roth; 5.14 in TSM)

No stable matching mechanism exists that makes it a dominant strategy for all hospitals to state their true preferences.

- Even when hospital-oriented GS (HGS) is executed and preferences are strict, some hospitals may benefit from misreporting their preferences.

Hospitals / Residents problem with Ties (**HRT**)

- In practice, residents' preference lists are short
- Hospitals' lists are generally long, so *ties* may be used – *Hospitals / Residents problem with Ties (HRT)*
- A hospital may be *indifferent* among several residents
 - E.g., $h_1: (r_1 \ r_3) \ r_2 \ (r_5 \ r_6 \ r_8)$
- An instance of **HRT** may not admit a hospital-optimal and/or a resident-optimal matching.
- A matching M is stable in an HRT instance I if and only if M is stable in some instance I' of HR obtained from I by breaking the ties [Manlove et al, 1999].

Couples in HR

- Pairs of residents who wish to be matched to geographically close hospitals form *couples*
- Each couple (r_i, r_j) ranks in order of preference a set of pairs of hospitals (h_p, h_q) representing the assignment of r_i to h_p and r_j to h_q
- Stability definition may be extended to this case [Roth, 1984; McDermid and Manlove, 2010; Biró et al, 2011]
- Gives the *H*ospitals / *R*esidents problem with *C*ouples (**HRC**)
- A stable matching need not exist:

$$\begin{array}{lcl}
 (r_1, r_2) : & (h_1, h_2) \\
 r_3 : & h_1 \quad h_2
 \end{array}$$

$$\begin{array}{lcl}
 h_1 : 1 : & r_1 \quad r_3 \quad r_2 \\
 h_2 : 1 : & r_1 \quad r_3 \quad r_2
 \end{array}$$

Couples in HR

- Pairs of residents who wish to be matched to geographically close hospitals form *couples*
- Each couple (r_i, r_j) ranks in order of preference a set of pairs of hospitals (h_p, h_q) representing the assignment of r_i to h_p and r_j to h_q
- Stability definition may be extended to this case [Roth, 1984; McDermid and Manlove, 2010; Biró et al, 2011]
- Gives the *H*ospitals / *R*esidents problem with *C*ouples (**HRC**)
- A stable matching need not exist:

$$\begin{array}{lcl}
 (r_1, r_2) : & (h_1, h_2) \\
 r_3 : & \textcircled{h_1} \textcircled{h_2}
 \end{array}$$

$$\begin{array}{lcl}
 h_1 : 1 : & r_1 \textcircled{r_3} r_2 \\
 h_2 : 1 : & r_1 \textcircled{r_3} r_2
 \end{array}$$

Couples in HR

- Pairs of residents who wish to be matched to geographically close hospitals form *couples*
- Each couple (r_i, r_j) ranks in order of preference a set of pairs of hospitals (h_p, h_q) representing the assignment of r_i to h_p and r_j to h_q
- Stability definition may be extended to this case [Roth, 1984; McDermid and Manlove, 2010; Biró et al, 2011]
- Gives the *H*ospitals / *R*esidents problem with *C*ouples (**HRC**)
- A stable matching need not exist:

$$\begin{array}{ll}
 (r_1, r_2) : & (h_1, h_2) \\
 r_3 : & h_1 \text{ } \star h_2
 \end{array}$$

$$\begin{array}{ll}
 h_1 : 1 : & r_1 \text{ } r_3 \text{ } r_2 \\
 h_2 : 1 : & r_1 \text{ } \star r_3 \text{ } r_2
 \end{array}$$

- Stable matchings can have different sizes

HRC is hard

- The problem of determining whether a stable matching exists in a given HRC instance is NP-complete, even if each hospital has capacity **1** and:
 - there are no single residents
[Ng and Hirschberg, 1988; Ronn, 1990]
 - there are no single residents, *and*
 - each couple has a preference list of length ≤ 2 , *and*
 - each hospital has a preference list of length ≤ 3
[Manlove and McBride, 2013]
- the preference list of each single resident, couple and hospital is derived from a strictly ordered master list of hospitals, pairs of hospitals and residents respectively [Biró et al, 2011], *and*
- each preference list is of length ≤ 3 , *and*
- the instance forms a “dual market”
[Manlove and McBride, 2013]

Algorithm for HRC


Not in the Exam

-
- Algorithm C described in [\[Biró et al, 2011\]](#):
 - A Gale-Shapley like heuristic
 - An *agent* is a single resident or a couple
 - Agents apply to entries on their preference lists
 - When a member of an assigned couple is rejected their partner must withdraw from their assigned hospital
 - This creates a vacancy – so any resident previously rejected by the hospital in question may have to be reconsidered
 - The *algorithm need not terminate*
 - if it terminates, the matching found is guaranteed to be stable (according to the definition of stability by [Biró et al, 2011](#))
 - it cannot terminate if there is no stable matching
 - it need not terminate even if there is a stable matching

Algorithm C: example

Not in the Exam

Resident preferences



Hospitals' preferences derived from the following master list:

$r_1 r_2 r_3 r_4 r_5 r_6 r_7 r_8$

Each hospital has capacity **1**

Stable matching

Not in the Exam

Resident preferences

\mathbf{r}_3 : \mathbf{h}_1 \mathbf{h}_5

\mathbf{r}_7 : \mathbf{h}_6 \mathbf{h}_8

$(\mathbf{r}_1, \mathbf{r}_5)$: $(\mathbf{h}_1, \mathbf{h}_2)$ $(\mathbf{h}_3, \mathbf{h}_6)$

$(\mathbf{r}_2, \mathbf{r}_4)$: $(\mathbf{h}_4, \mathbf{h}_5)$ $(\mathbf{h}_1, \mathbf{h}_2)$ $(\mathbf{h}_3, \mathbf{h}_7)$

$(\mathbf{r}_6, \mathbf{r}_8)$: $(\mathbf{h}_6, \mathbf{h}_8)$

Hospitals' preferences

\mathbf{r}_1 \mathbf{r}_2 \mathbf{r}_3 \mathbf{r}_4 \mathbf{r}_5 \mathbf{r}_6 \mathbf{r}_7 \mathbf{r}_8

Each hospital has capacity **1**

Stable matching: $M = \{(r_1, h_3), (r_2, h_1), (r_3, h_5), (r_4, h_2), (r_5, h_6), (r_7, h_8)\}$

Empirical evaluation

Not in the Exam

-
- Extensive empirical evaluation due to [\[Biró et al, 2011\]](#):
 - Compared 5 variants of Algorithm C against 10 other algorithms
 - Instances generated with varying:
 - sizes
 - numbers of couples
 - densities of the “compatibility matrix”
 - lengths of time given to each instance
 - Measured proportion of instances found to admit a stable matching
 - Clear conclusion:
 - high likelihood of finding a stable matching (with Algorithm C) if the number / proportion of couples is low

What was in this lecture

- Hospitals/Residents problem (HR) and HR with ties (HRT)
- Resident-oriented and Hospital-oriented GS for HR
- Rural hospitals theorem
- DS truthfulness in HR
- Couples in HR

Acknowledgement

Some of the slides in this lecture were based on the slides by **David Manlove**.

Books

- **Algorithmics of Matching under Preferences** by David F. Manlove.
- **Two-Sided Matching: A Study in Game-Theoretic Modeling and Analysis (TSM)** by Alvin E. Roth, Marilda A. Oliviera Sotomayor.
- **Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations (MAS)** by Yoav Shoham and Kevin Leyton-Brown
- **Algorithmic Game Theory (AGT)**, edited by Noam Nisan, Tim Roughgarden, Eva Tardos, Vijay V. Vazirani

Optional additional reading

- On the “history of labor market for medical interns”
 - **TSM** section 1.1, or alternatively
 - *The Evolution of the Labor Market for Medical Interns and Residents: A Case Study in Game Theory* by Alvin E. Roth, Journal of Political Economy, Vol 92, 1984
- The NIMP algorithm
 - **TSM** section 5.4
 - *The Evolution of the Labor Market for Medical Interns and Residents: A Case Study in Game Theory* by Alvin E. Roth, Journal of Political Economy, Vol 92, 1984