

COMP 3225

Natural Language Processing

Sequence Processing, Transformer and
Attention

Stuart E. Middleton

sem03@soton.ac.uk

University of Southampton

Copyright University of Southampton 2021.

Content for internal use at University of Southampton only.

Slides may include content publicly shared for education purposes via <https://web.stanford.edu/~jurafsky/slp3/>

Overview

- LSTM
- GRU
- Gated Units, Layers and Networks
- Transformer

- <break - discussion point>

- Self-attention
- Transformer for Text Completion Task


LSTM

- Long Short-Term Memory (LSTM)
- Long-distant information is critical to many language apps
 - Information encoded by a single RNN layer tends to be fairly local
 - Information is lost each training step, so encoded memory of tokens far away in sequence degrades pretty quickly
- Vanishing gradients problem
 - Gradients get smaller and smaller as Gradient Descent progresses to deeper stacked layers - they 'vanish'
 - Connection weights are virtually unchanged, so training loss does not converge
- Exploding gradients problem
 - The opposite can also happen, where gradients get larger and larger and the training loss again does not converge
 - In general these problems lead to unstable gradients

LSTM


- Long short-term memory (LSTM)

- Forgets information no longer relevant (**forget gate**)
- Adds new information (**input gate** sometimes called **add gate**)
- Gate** = layers \rightarrow logistic activation function σ (value 0..1), followed by a pointwise multiplication to provide a type of binary mark




GRU

- Gated Recurrent Unit (GRU)
 - Merges both long and short term state vectors
 - Simpler (less layers) and often performs as well as LSTM




Gated Units, Layers and Networks


- Stacked RNN, LSTM and GRU layers can be experimented with easily using frameworks like Tensorflow and PyTorch




MLP cell



RNN cell



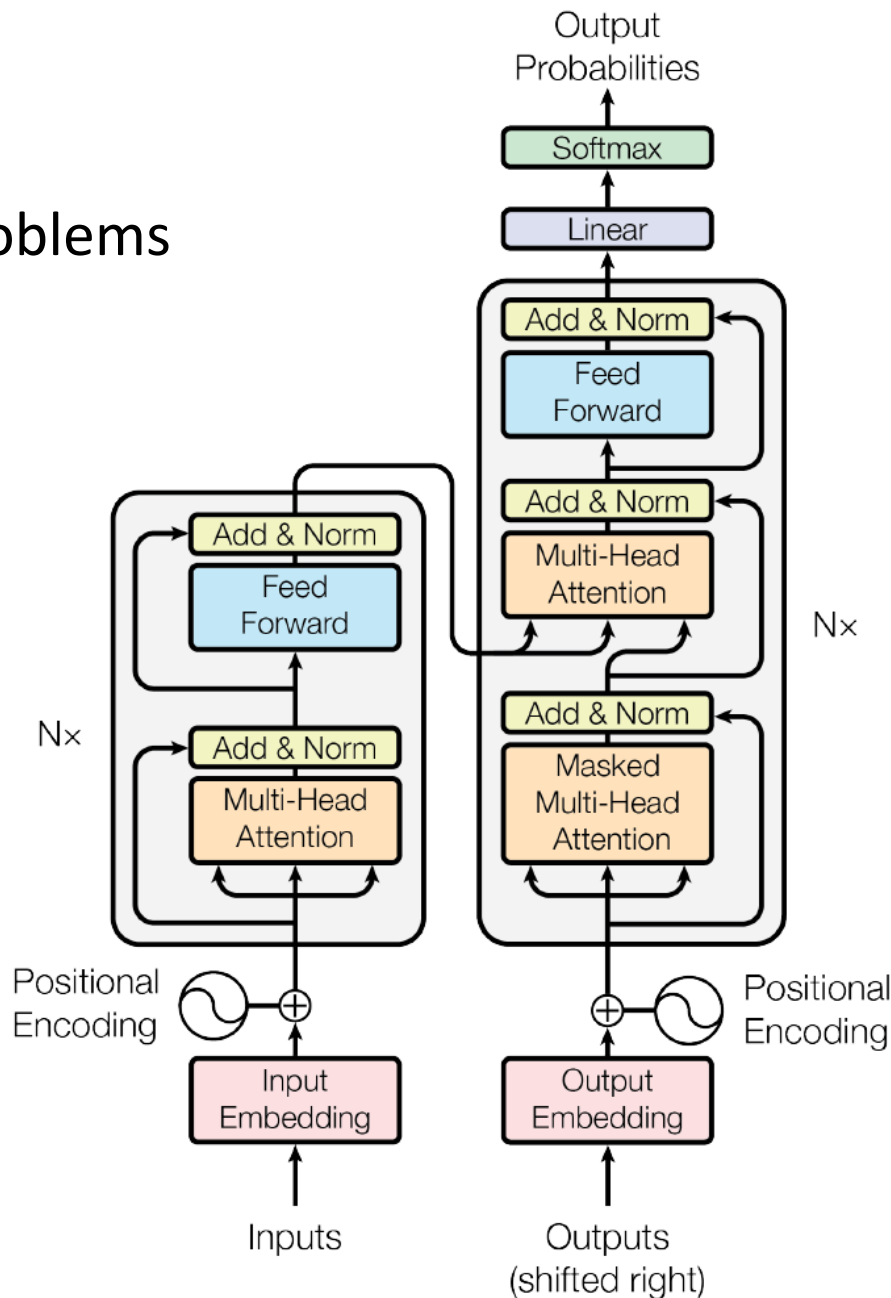
LSTM cell



GRU cell


Transformer

- Even LSTM and GRU's have problems with very long sequences
- **Transformer architecture**
 - Uses attention layers not sequential RNN layers
 - This architecture has led to many advances in NLP model performance



Self-attention

- Attention layers >> compare item to other items to reveal their relevance in the current context
 - Self-attention compares to other elements within same sequence
- Self-attention layer maps input sequences (x_1, \dots, x_n) to output sequences of the same length (y_1, \dots, y_n)
 - Layer computes for position i the y_i using input $x_1 \dots x_i$ (uses context)
 - Layer computation independent of other layers (parallelizable)



Self-attention

- Compare items (embedding vector for a word) in a sequence using dot product using a score function

$$\text{score}(x_i, x_j) = x_i \cdot x_j$$

$$\alpha_{ij} = \text{softmax}(\text{score}(x_i, x_j)) \quad \forall j \leq i \quad \text{j = all values up to i}$$

$$y_i = \sum_{j \leq i} \alpha_{ij} x_j$$

- For seq $x_1..x_3$ we need compute $\text{score}(x_3, x_1)$, $\text{score}(x_3, x_2)$, $\text{score}(x_3, x_3)$
- Normalize scores using softmax to create a vector of attention weights α
- Compute output by summing inputs in seq $x_1..x_i$ weighted by α
- But ... to allow learning we need trainable weights

Break

- Panopto Quiz - discussion point
- How can we provide trainable weights for the attention function?

We don't need to! we have the attention weights α already

Use a LSTM layer

Use an weight vector

Use an weight matrix

$$y_i = \sum_{j \leq i} \alpha_{ij} x_j$$

Break

- Panopto Quiz - discussion point
- How can we provide trainable weights for the attention function?

We don't need to! we have the attention weights α already >> we need weights to train

Use a LSTM layer >> there is no sequence structure in an attention layer

Use an weight vector >> No

Use an weight matrix >> Yes, actually several matrices as you will see next

$$y_i = \sum_{j \leq i} \alpha_{ij} x_j$$

Self-attention

- Imagine encoder learns a dictionary
 - concept \rightarrow lexical term == key \rightarrow query
'They played chess' \gg 'subject' \rightarrow 'They'; 'verb' \rightarrow 'played'
 - Concepts are learnt in latent space do not map directly to single words
- Decoder can 'lookup' values for concepts it think go next in seq
- Transformer uses embeddings to provide this dictionary which is then differentiable (for gradient decent)
 - Query Q = matrix $[n_{\text{queries}}, d_{\text{keys}}]$, where n_{queries} = input seq length
 - Key K = matrix $[n_{\text{keys}}, d_{\text{keys}}]$, where n_{keys} = number of keys
 - Value V = matrix $[n_{\text{keys}}, d_{\text{values}}]$
 - d_{keys} = dimensions of query and key embedding

$$q_i = W^Q x_i; \quad k_i = W^K x_i; \quad v_i = W^V x_i$$

$$\text{score}(x_i, x_j) = \frac{q_i \cdot k_j}{\sqrt{d_k}}$$

$$y_i = \sum_{j \leq i} \alpha_{ij} v_j$$

**layer output =
weighted sum of (attention x value)
for each word in seq**

Self-attention

- Since computation of output q_i can be done independently, we can use matrix multiplication to parallelize training

$$q_i = W^Q x_i; \quad k_i = W^K x_i; \quad v_i = W^V x_i$$

$$Q = W^Q X; \quad K = W^K X; \quad V = W^V X$$

$$\text{SelfAttention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

- For language modeling we do not want to use sequence information after the current position (that would be cheating)
- We can blank (i.e. give a value of $-\infty$) matrix values after the current position to hide them

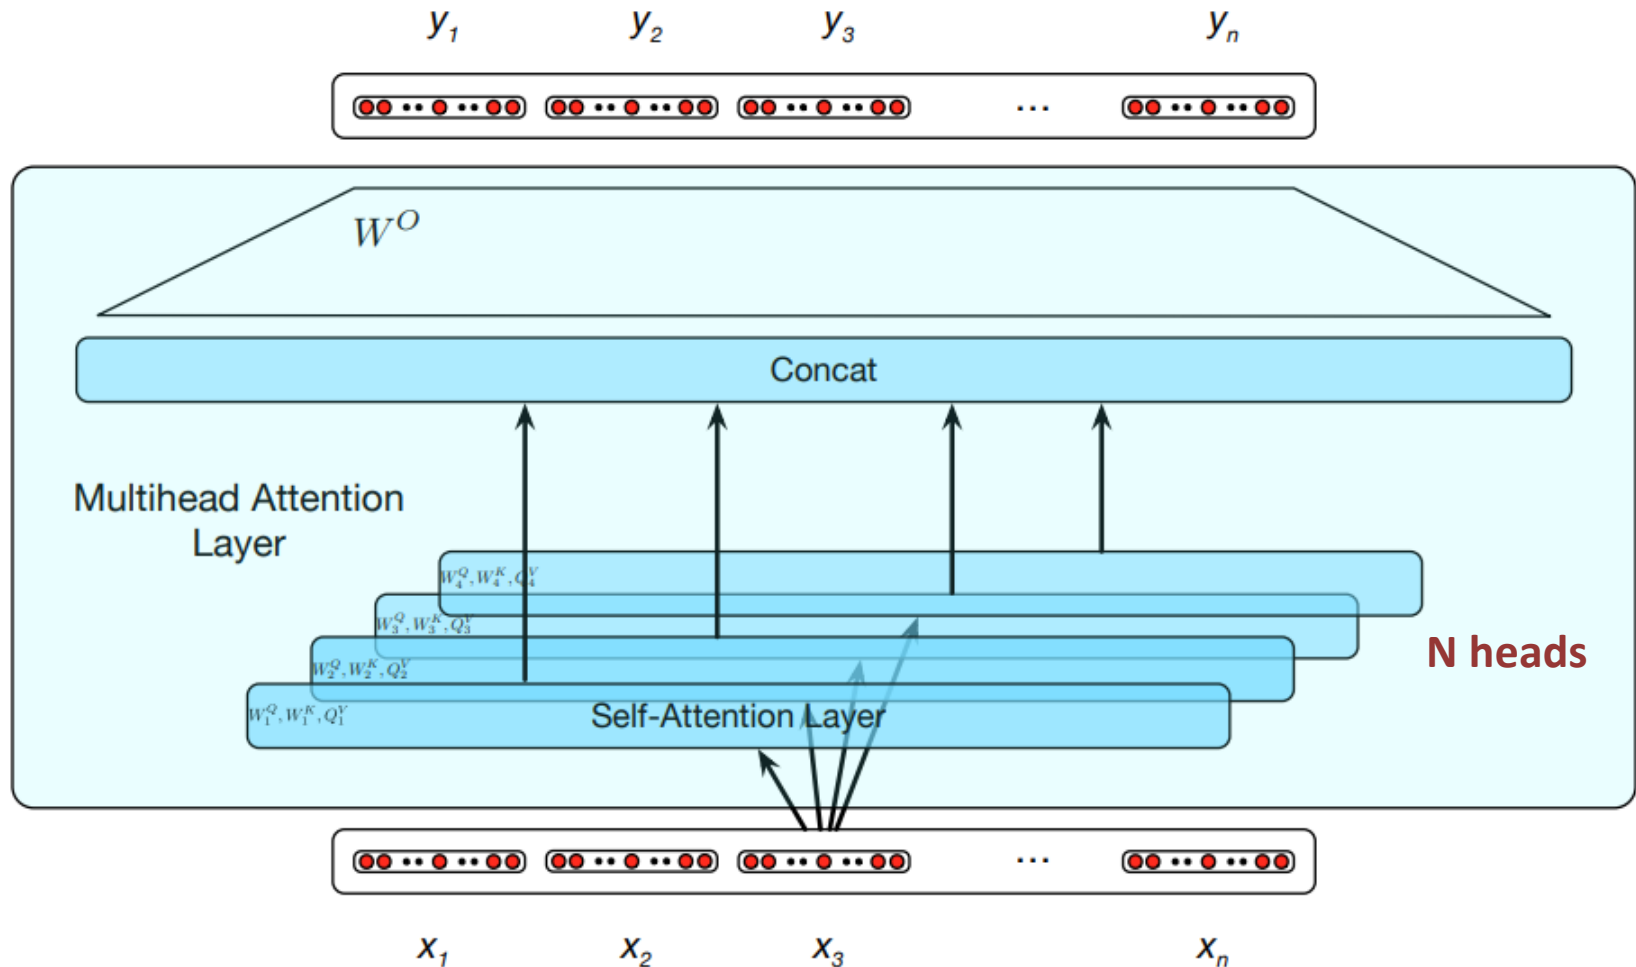
Self-attention

- Positional embeddings

- The position of words in sequences is important for language problems
- Transformer attention layers do not encode the position of words (unlike RNNs), so it uses an additional positional embedding
- These can be learned during training (e.g. additional position input alongside words) or generated using a positional encoding function (e.g. cosine/sine as per Vaswani 2017 paper)
- Positional embeddings have the same dimensions as word embeddings, so are just summed together to give a position-aware input embedding


Self-attention

- Multi-head Self-attention Layers
 - A self-attention layer is called a head
 - There can be N heads in a model providing a deep learning stack



Transformer for Text Completion Task

- Example using Transformer model for text completion
- Input = sequence of words
- Output = prediction of words to complete sequence
- Transformer architectures have become a 'go to' model for NLP



Required Reading

- LSTM, GRU, Transformer and Self-attention
 - Jurafsky and Martin, Speech and Language Processing, 3rd edition (online)
>> chapter 9
- Transformer and Self-attention
 - Aurelien Geron, Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, O'Reilly , 2017
>> Chapter 16 'attention mechanisms'
- LSTM, GRU (optional)
 - Aurelien Geron, Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, O'Reilly , 2017
>> Chapter 15 'handling long sequences'

Questions

- Panopto Quiz - 1 minute brainstorm for interactive questions
Please write down in Panopto quiz in **1 minute** two or three questions that you would like to have answered at the next interactive session.

Do it **right now** while its fresh.

Take a screen shot of your questions and **bring them with you** at the interactive session so you have something to ask.